

CS 354

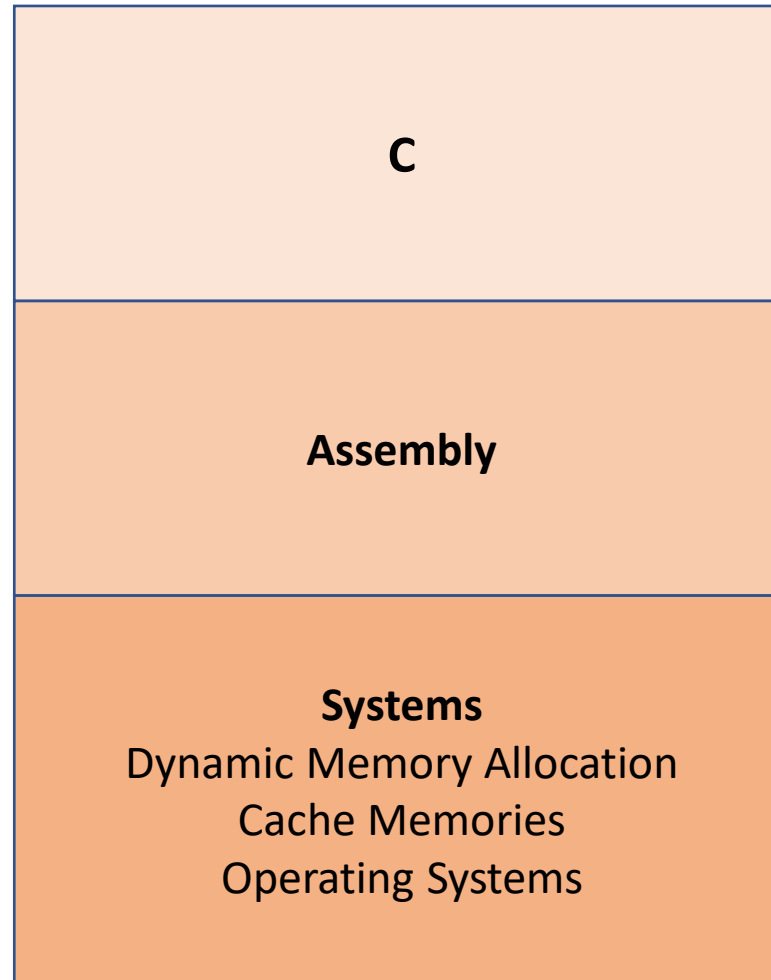
Machine Organization and Programming

Lecture 25

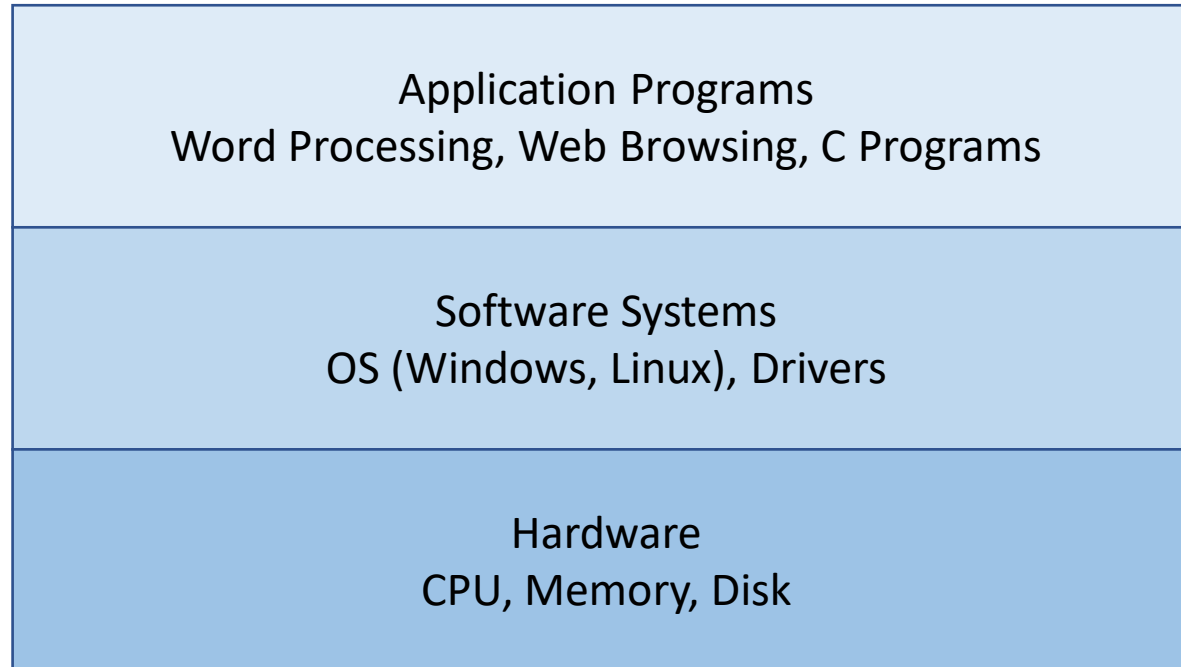
Michael Doescher
Summer 2020

Intro to Operating Systems

Course Overview



How Computers Work



What is an Operating System

- A program that makes your computer much easier to use
- A program that manages computer resources and provides common services needed by applications
- Input / Output, Memory Management, Time Sharing
- Takes care of the details to allow your computer to run two programs at the 'same time'

How does a computer run two programs at the same time?

- If we only have 1 CPU?
- Memory Address Space is a contiguous block from addr: 0 to addr: MAX

What does the CPU do?

sum.c

.....

a = 1;

b = 2;

c = a + b;

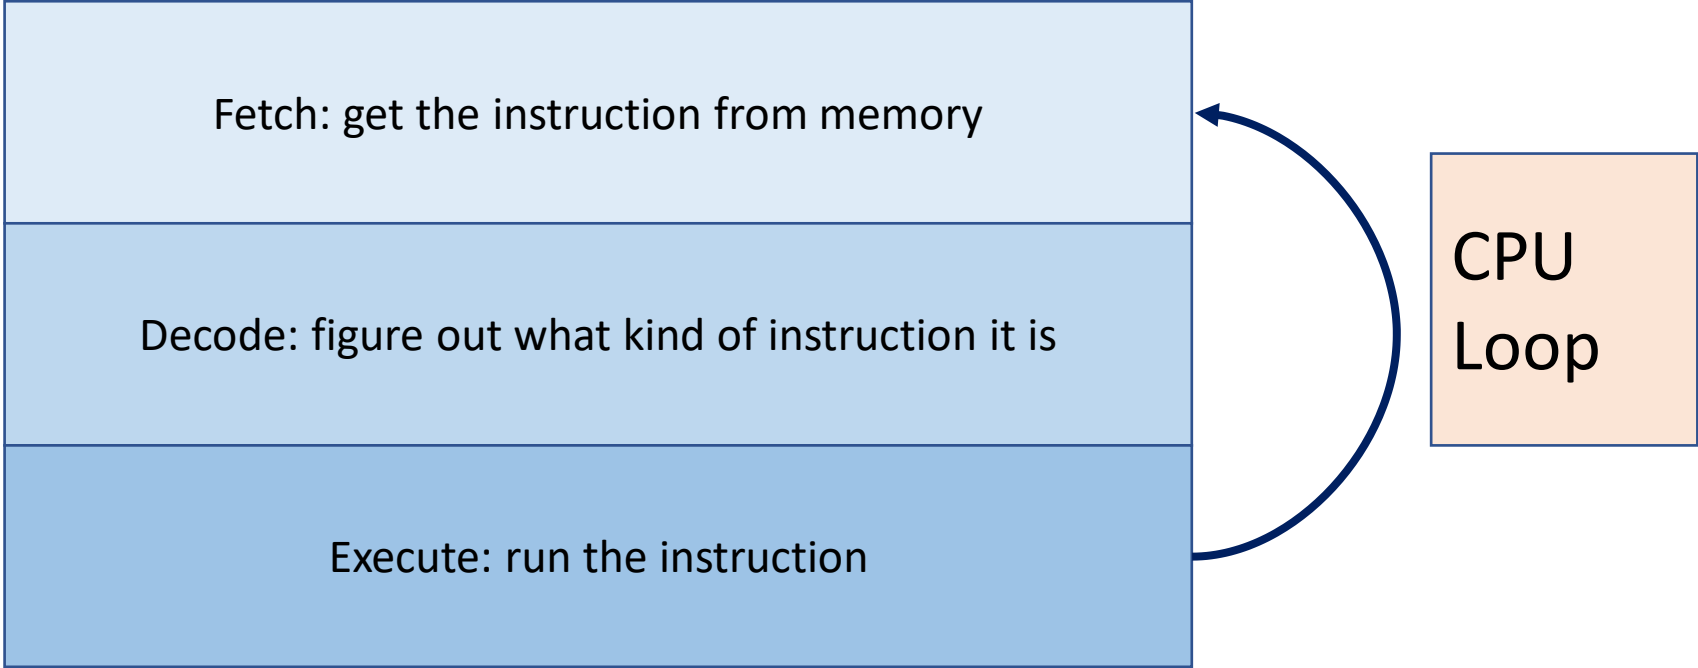
.....

Fetch: get the instruction from memory

Decode: figure out what kind of instruction it is

Execute: run the instruction

CPU
Loop




The diagram illustrates the CPU Loop process. On the left, a light green box contains the source code for 'sum.c', showing variable declarations and an addition operation. To the right, a large light blue box is divided into three horizontal sections representing the CPU's internal steps: 'Fetch: get the instruction from memory', 'Decode: figure out what kind of instruction it is', and 'Execute: run the instruction'. To the right of this box is a light orange box labeled 'CPU Loop'. A curved arrow originates from the bottom of the 'Execute' section and points back to the top of the 'Fetch' section, indicating a continuous loop of processing instructions.

Two Programs - One CPU

Time Sharing

Switch back and forth between processes

OS figures out how to schedule when each program runs on the CPU



Shared Address Space

Memory is a linear array of memory address from 0 to MAX

Programs have access to the entire address space.

What if Program A writes to addr: 0xFA0
Then Program B writes to addr: 0xFA0

Virtual address space

```
main:
0x100 | instr 1
0x104 | instr 2
0x108 | call func
0x10C | instr 3
```

```
func:
0x204 | instr 41
0x208 | instr 42
0x20C | ret
0x201 | instr 44
```

```
Heap:
0x410 | 0x18
```

```
stack:
0xFA0 | 0x03
0xFA4 | 0x10C
0xFA8 | 0x29
```

```
main:
0x100 | instr 1
0x104 | instr 2
0x108 | call func
0x10C | instr 3
```

```
func:
0x204 | instr 41
0x208 | instr 42
0x20C | ret
0x201 | instr 44
```

```
Heap:
0x410 | 0x18
```

```
stack:
0xFA0 | 0x03
0xFA4 | 0x10C
0xFA8 | 0x29
```


Shared Address Space

Virtual address space

Addresses used by programs are not the addresses used in physical memory

OS lets us treat the memory that a program sees as the physical memory

OS keeps track of mapping where memory used by each program is really stored in physical memory

```
main:
0x100 | instr 1
0x104 | instr 2
0x108 | call func
0x10C | instr 3
```

```
func:
0x204 | instr 41
0x208 | instr 42
0x20C | ret
0x201 | instr 44
```

```
Heap:
0x410 | 0x18
```

```
stack:
0xFA0 | 0x03
0xFA4 | 0x10C
0xFA8 | 0x29
```

```
main:
0x100 | instr 1
0x104 | instr 2
0x108 | call func
0x10C | instr 3
```

```
func:
0x204 | instr 41
0x208 | instr 42
0x20C | ret
0x201 | instr 44
```

```
Heap:
0x410 | 0x18
```

```
stack:
0xFA0 | 0x03
0xFA4 | 0x10C
0xFA8 | 0x29
```


CPU Virtualization

Program A runs on CPU A

Program B runs on CPU B


Each program sees the illusion that they have access to the entire CPU

Time sharing creates this illusion.




Processes


Program




Process



Process



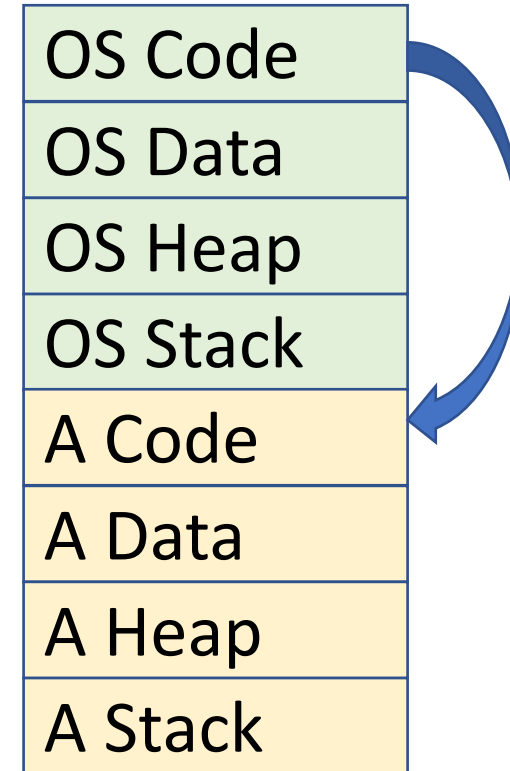

How do you switch processes?




Resources Used - State

- Registers
- Memory code, data, stack, heap (evict?)
- Address Space
- Open files

How do you switch processes?



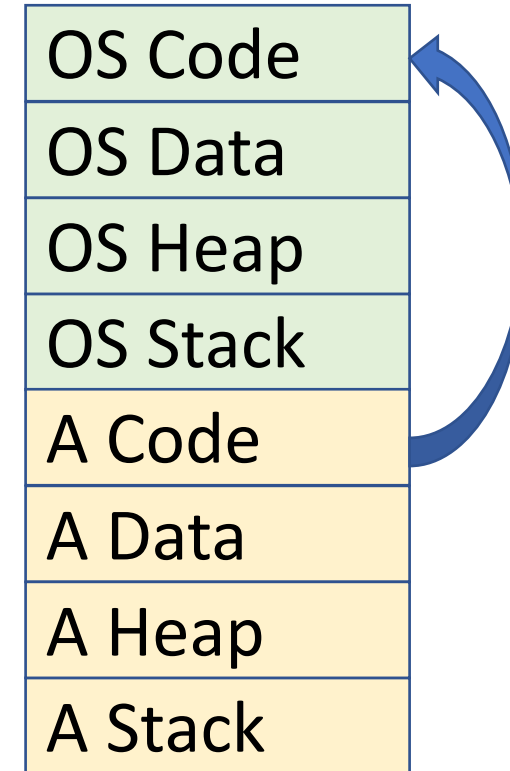
How do you switch back?




`yield();`

Requires all process to play nice

```
while(1) {  
    ....  
}
```



OS needs to regain control



```
main:
0x100 | instr 1
0x104 | instr 2
0x108 | call func
0x10C | instr 3

func:
0x204 | instr 41
0x208 | instr 42
0x20C | ret
0x201 | instr 44
```

Interrupt

Hardware support

Timer interrupt –
duration set by OS

Hardware Interrupt Handler

%eip
0xFFFF0

```
Interrupt Handler:
0xFFFF0 .....
0xFFFF4 .....
0xFFFF8 .....
0xFFFFC .....
```

CPU Virtualization Mechanisms

1. Timer Interrupt
- 2.

At Boot time (power on the computer)

OS tells the hardware the location of the interrupt handlers

Types of interrupts

- timer
- keyboard
- disk

OS sets up the timer interrupt - a physical clock that just produces signals at a regular interval

Problems

1. Timer Interrupt

2.

- Timer can be switched on and off
- OS can change the address of the timer interrupt code

OS doesn't want to be interrupted!

What if a user process also switches off the timer?

Solution

1. Timer Interrupt
2. **Privilege Mode**


set to 0 then user mode

set to 1 then **kernel** mode (the core of the OS)

- core = scheduler, memory virtualization, ...
- not core – device drivers, input/output, ...


When kernel mode is set – any instruction can be executed

Example: Turning off the timer interrupt is a privileged instruction

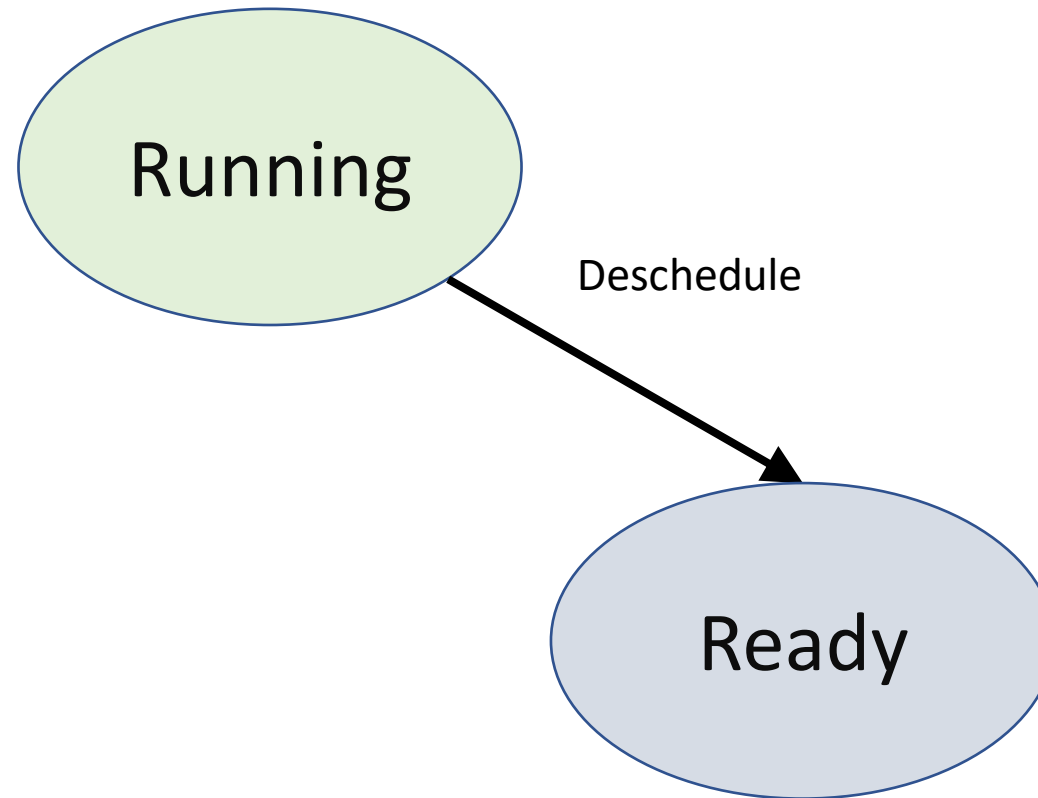


Status Register

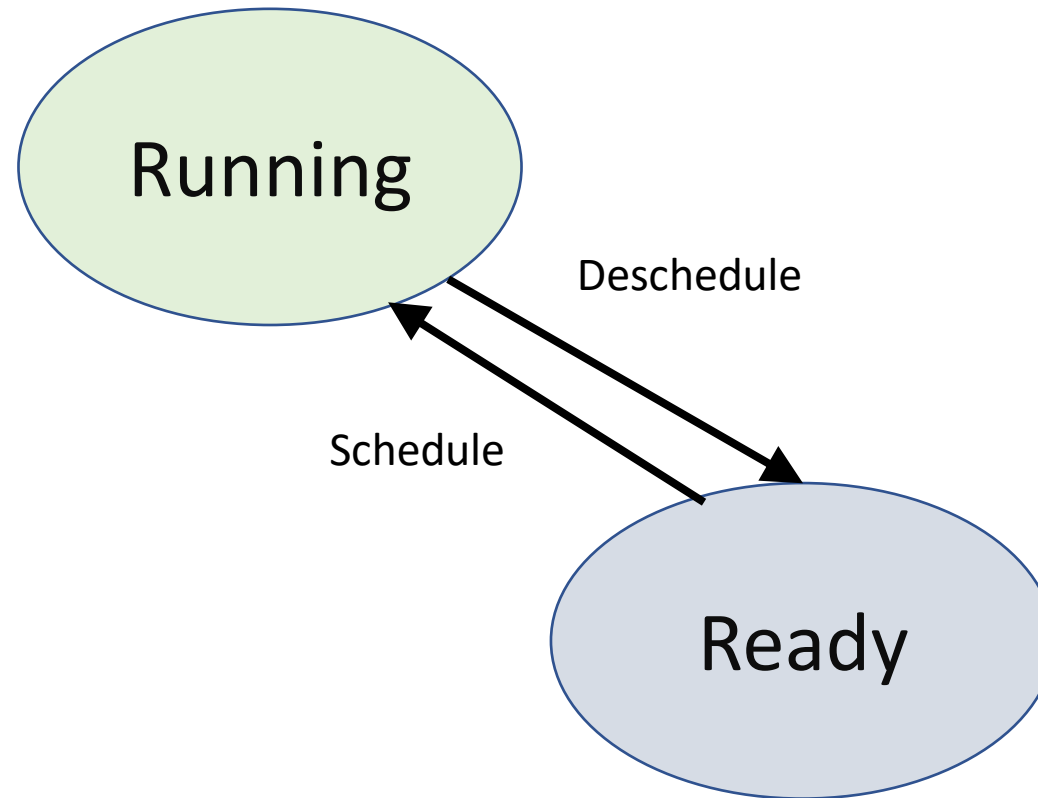
Process State



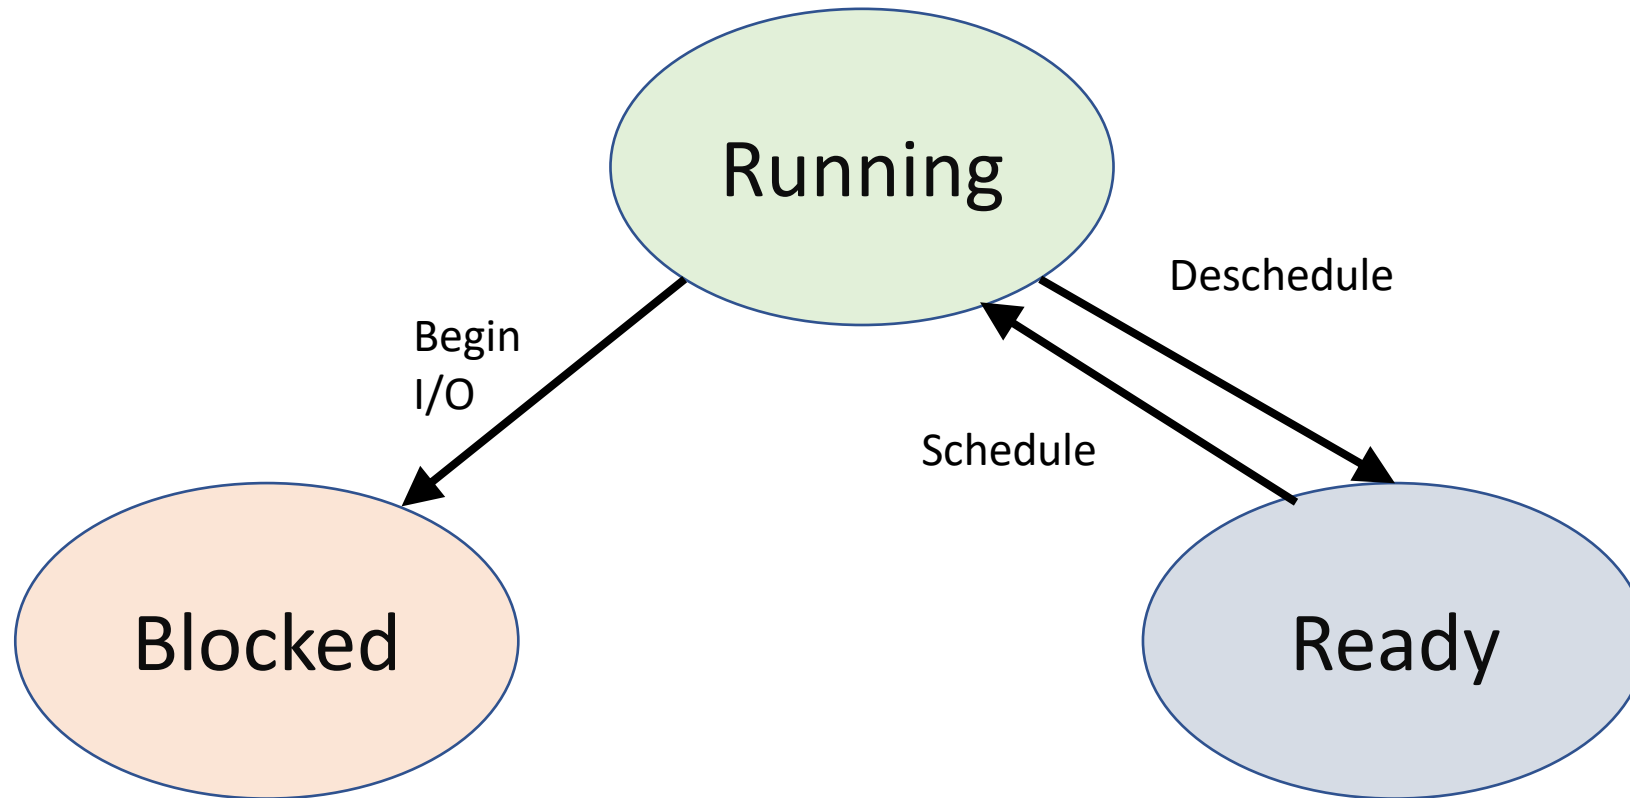
Process State



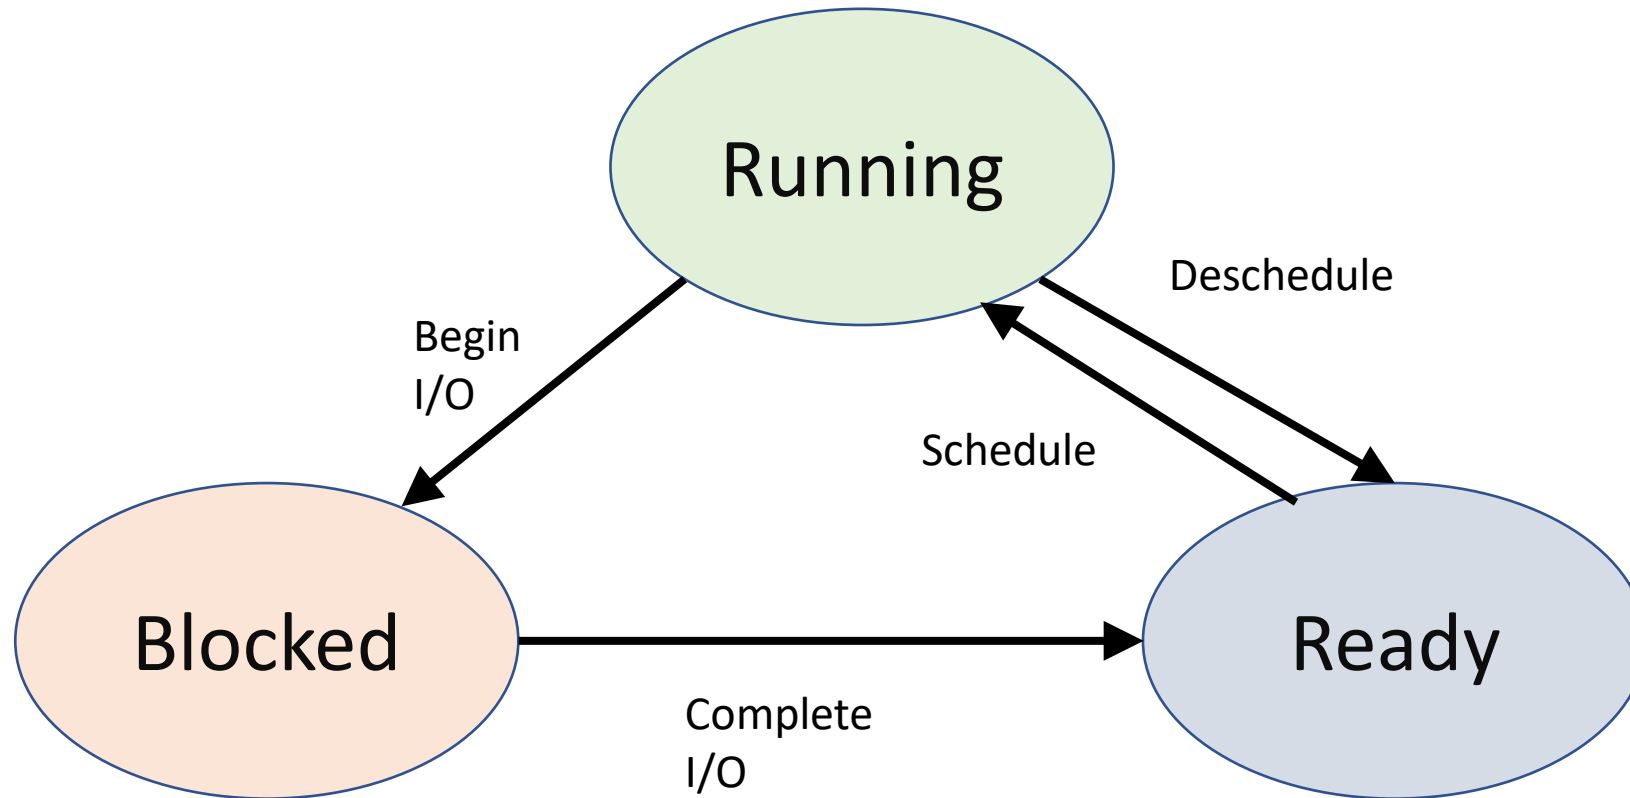
Process State




Process State



Process State



Process State



CS 354

Machine Organization and Programming

Lecture 25

Michael Doescher
Summer 2020

Intro to Operating Systems