

CS 354

Machine Organization and Programming

Michael Doescher
Spring 2021

Intro to Operating Systems II

- CPU Virtualization
- Memory Virtualization
- Concurrency / Threads



Account



Dashboard



Courses



Calendar



Inbox



History



Commons



Help

SP21 COMPSCI 354 003 > Pages > Week 13 Readings

Student View

Spring 2020-2021

View All Pages

Published

Edit



Home

Announcements

Modules



People

Grades

BBCollaborate Ultra

Kaltura My Media

Course Summary

Course Syllabus
(AEFIS)

Direct Evidence of
Student Learning
(AEFIS)

Assignments



Discussions



Week 13 Readings

- [Introduction to OS.pdf](#) ↓
- [Processes.pdf](#) ↓
- [Direct_Execution.pdf](#) ↓
- [Address_Space.pdf](#) ↓
- [Address_Translation.pdf](#) ↓
- [Paging.pdf](#) ↓
- [Threads.pdf](#) ↓

The Process

- 10000 processes at a time?
- With 1 CPU?
- “Virtualizing the CPU” emulating many CPU’s
- OS creates the illusion of an endless supply of CPUs by time sharing
- Hardware Resources required for the “context switching”
- OS Resources to make decisions about scheduling

State

- List of all machine resources used by a process
 - Memory (instructions, data, heap, stack) -> The Address Space
 - Registers
 - I/O
 - files
 - standard in, standard out, error stream
 - network connections

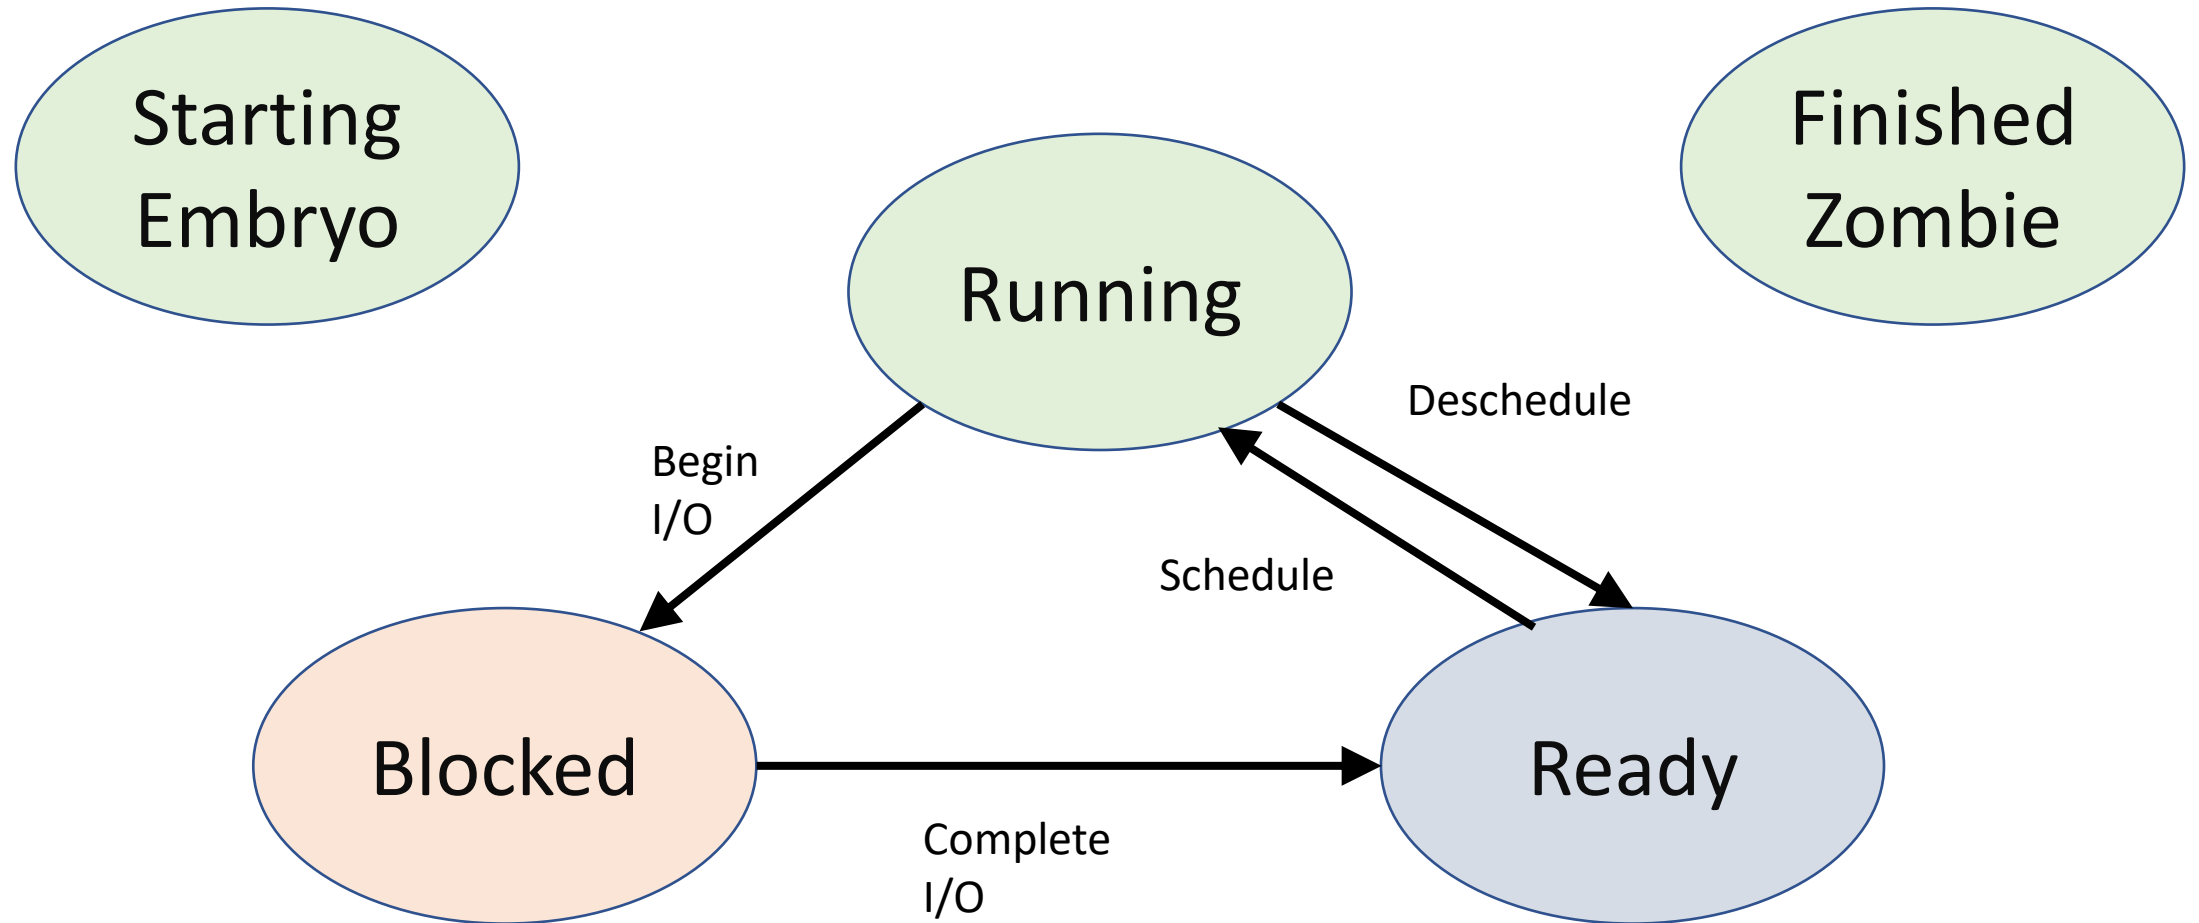
System Calls

- Create
- Destroy
- Wait
- Control
- Status

Create

- Add entry to the process list
- Copy Code and Data to Memory (load)
- Allocate Memory for Heap and Stack
- Push Command Line Parameters onto the Stack
- Setup I/O – standard in , standard out, error
- Clear Registers
- Transfer Control to 1st instruction of main

Process State



Direct Execution

- OS Creates Process
- Transfer Control to Process
 - Run Main
 - Return 0
- Destroy
- Free Process Memory
- Remove from Process List

Issues

- No Limitations
- Infinite Loop
- Security
 - restricted operations

Limited Direct Execution

- Time – Sharing
- User Mode and Kernel Mode
- Trap Instructions + Return from Trap
- Trap and Interrupt Tables set up at boot time
 - Where is the code to run each system call
 - Address of the SysCall Handler is stored in a register at boot time

System Call Protocol

1) OS

- Create Process
- Push Registers to Kernel Stack
- Push PC to Kernel Stack
- Return from Trap

2) Hardware

- Restore Registers (from Kernel Stack)
- User Mode
- Restore PC (from Kernel Stack)

3) Process

- Run Main
- System Call – Trap to OS
- With System Call Number

4) Hardware

- Push Registers to Kernel Stack
- Kernel Mode
- Jump to Trap Handler (PC)

5) OS

- Handle Trap
- Return from Trap

6) Hardware

- Go to Step 2

7) OS

- If Trap is exit / return from main
- Free memory
- Remove from Process List

CS 354

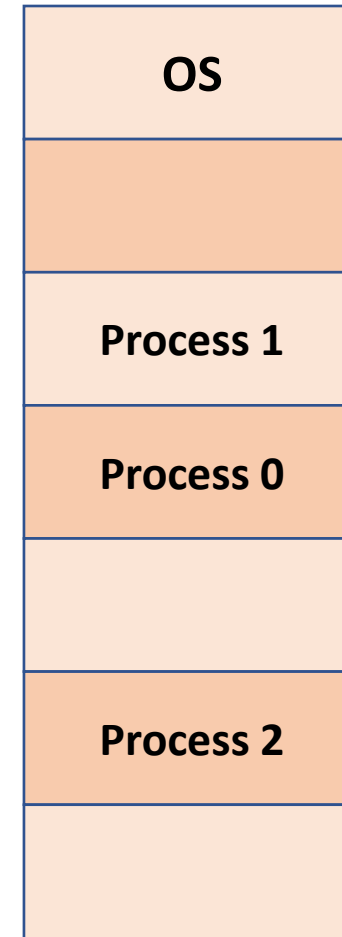
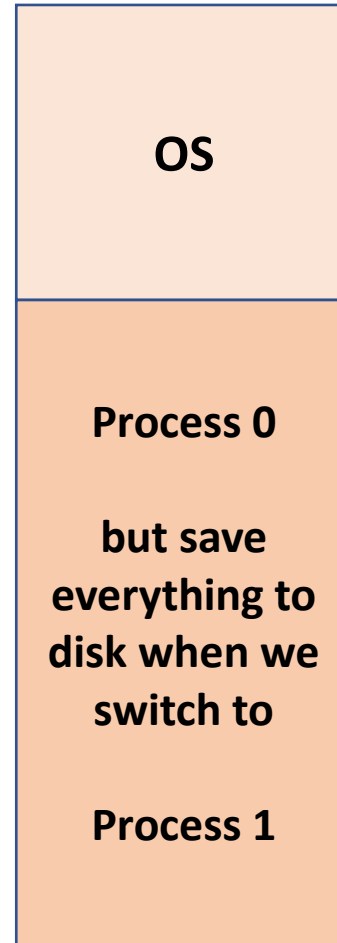
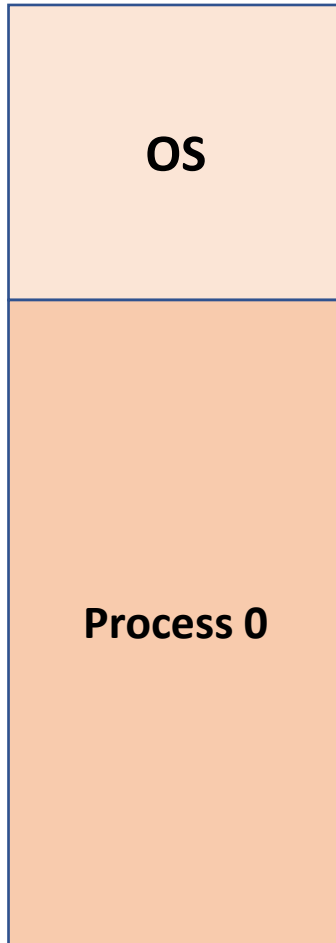
Machine Organization and Programming

Michael Doescher
Spring 2021

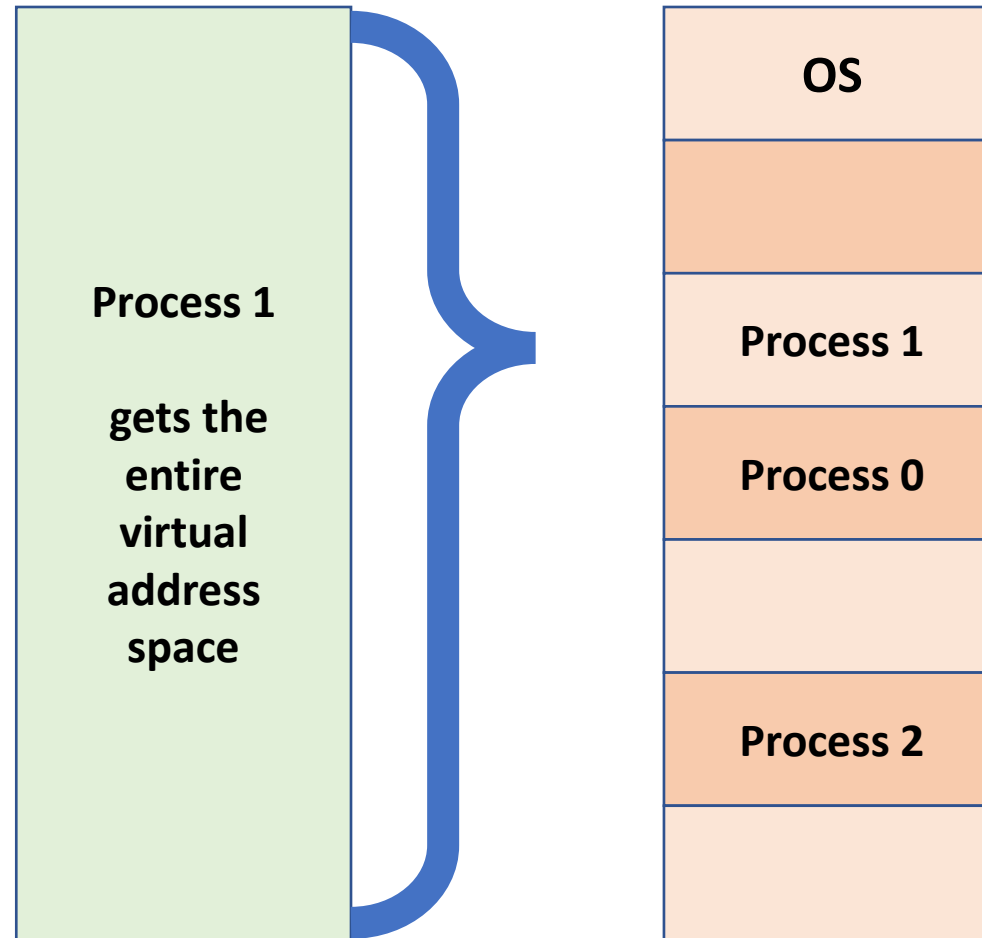
Intro to Operating Systems II

- CPU Virtualization
- **Memory Virtualization**
- Concurrency / Threads

Address Space Models



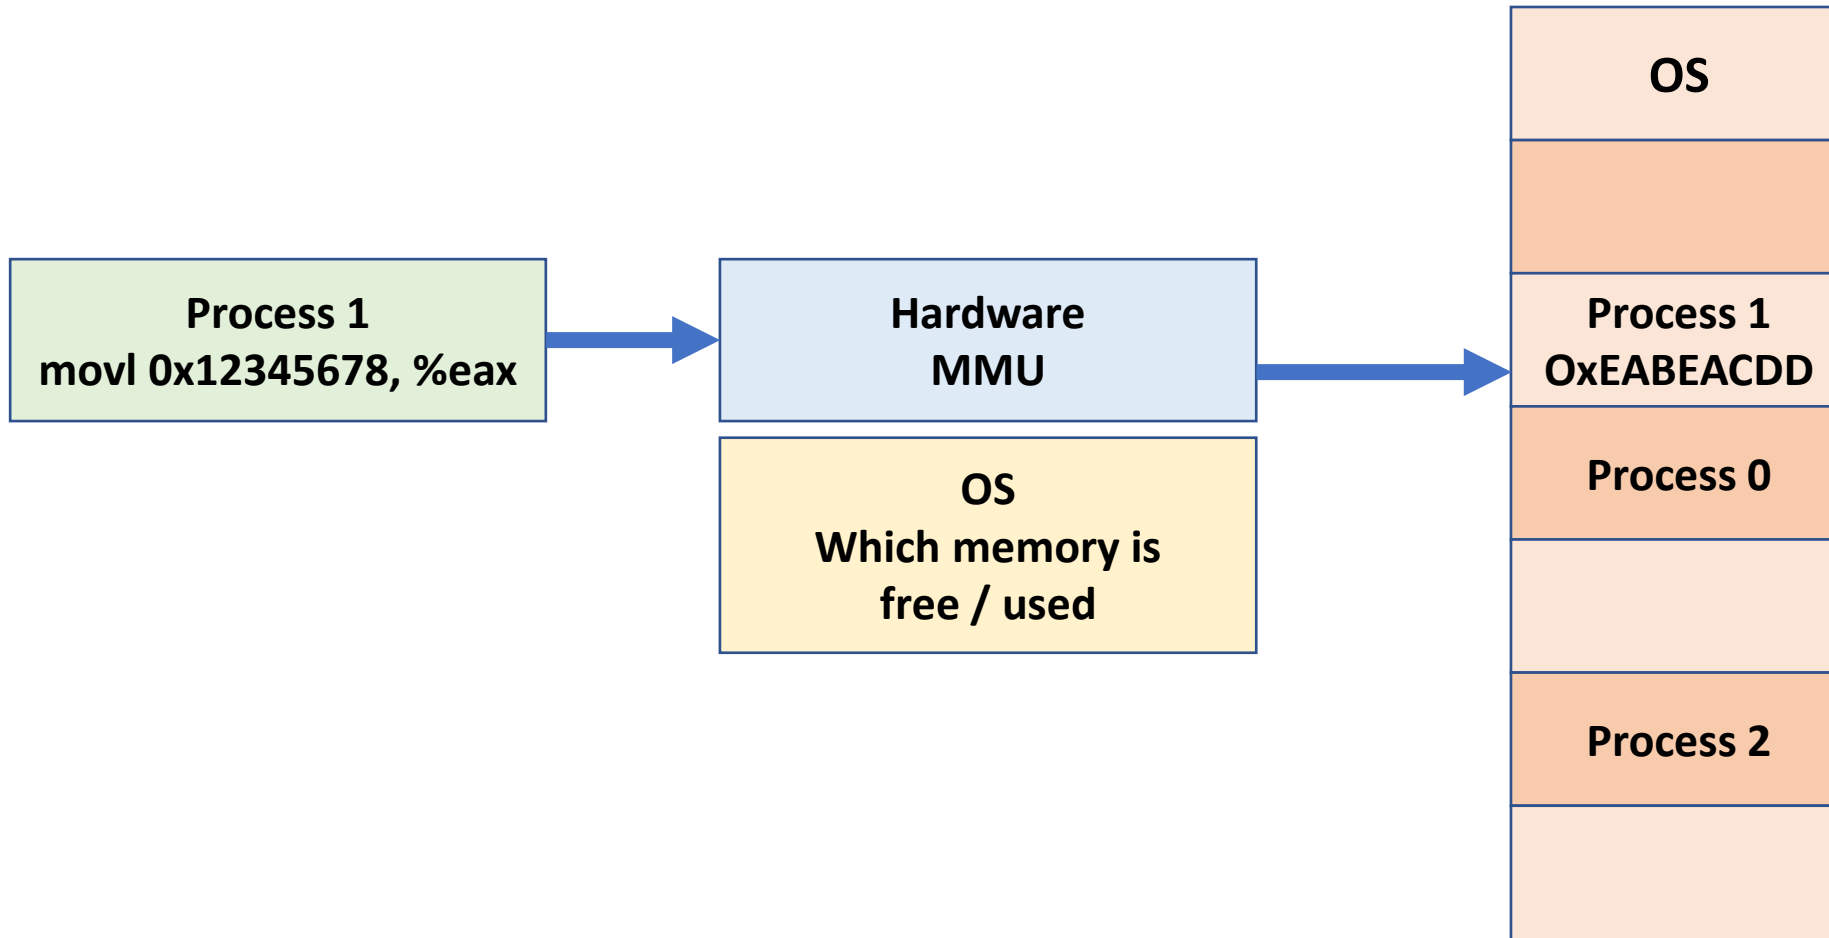
Address Space Virtualization



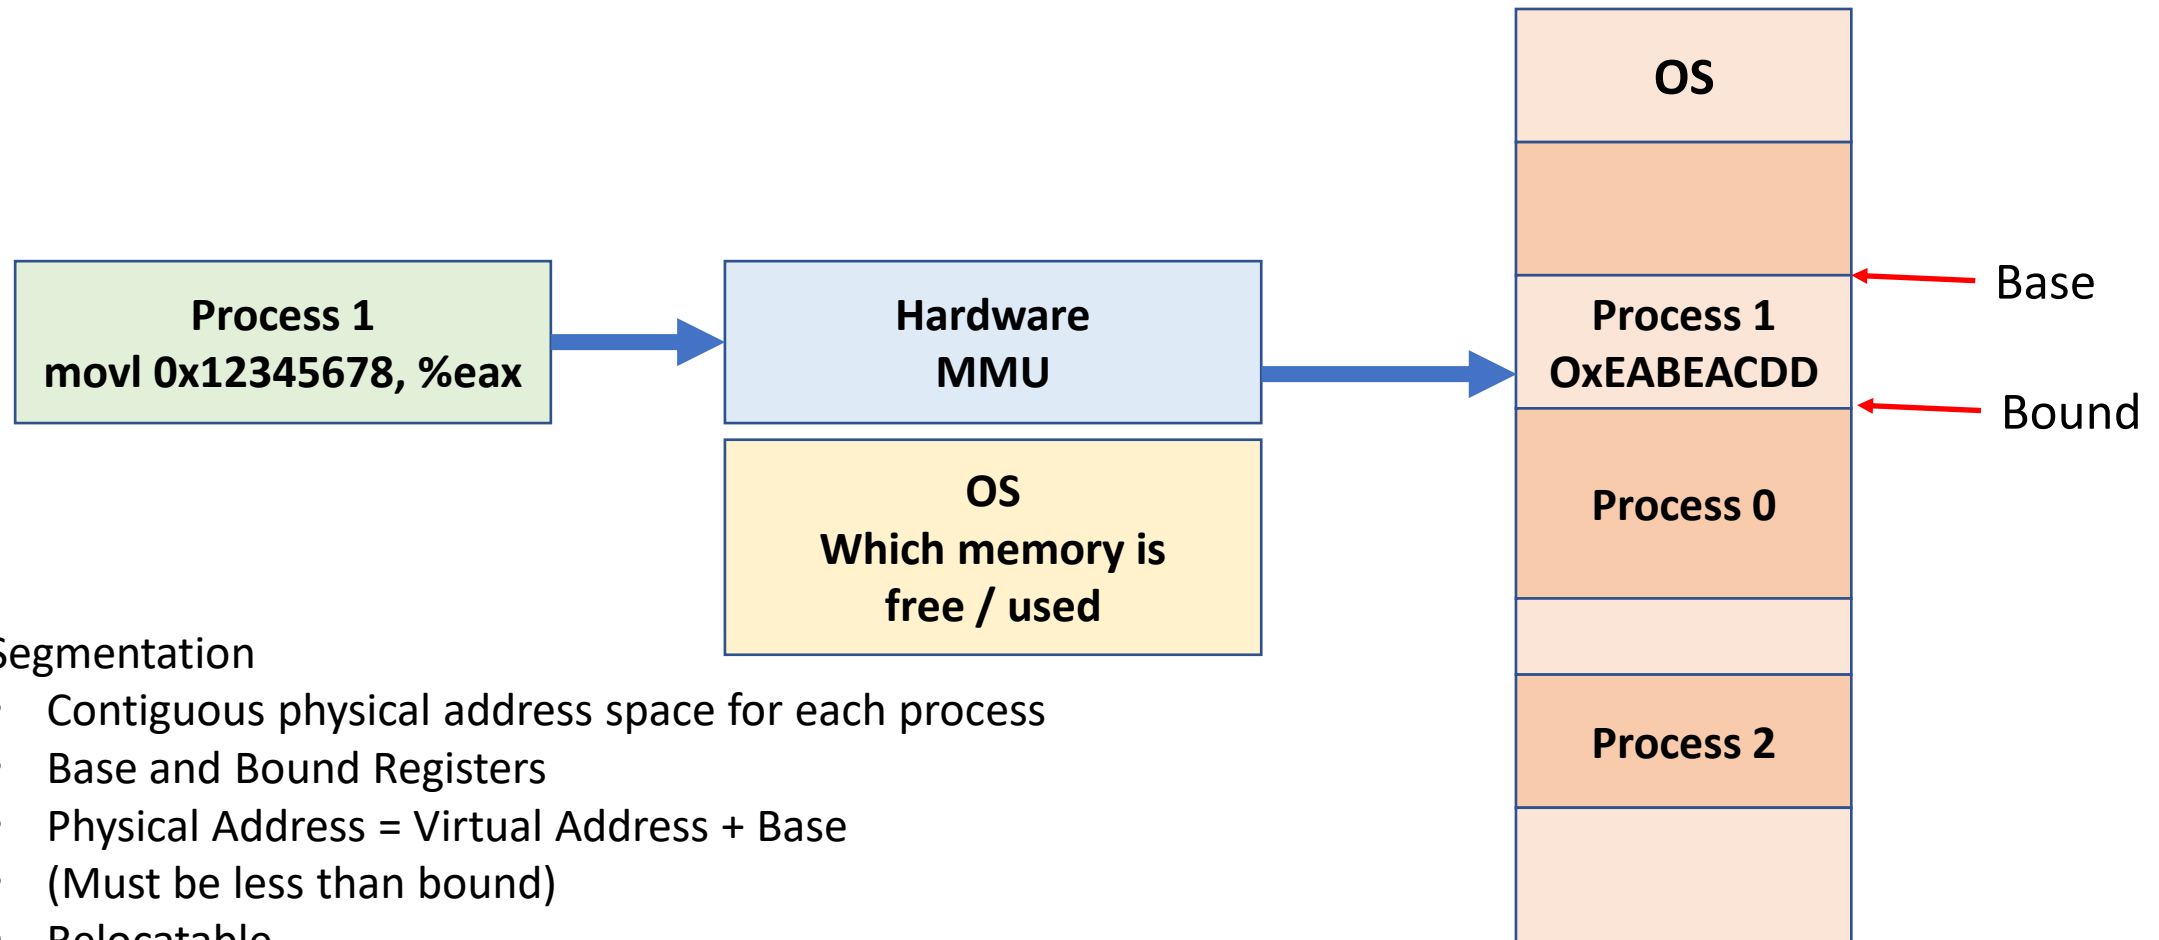
Goals

- Hide all details from the user / programmer
- Efficiency
- Protection / Isolation
- Allocation Policies – OS takes care of this
- Easy to Program

Address Space Virtualization



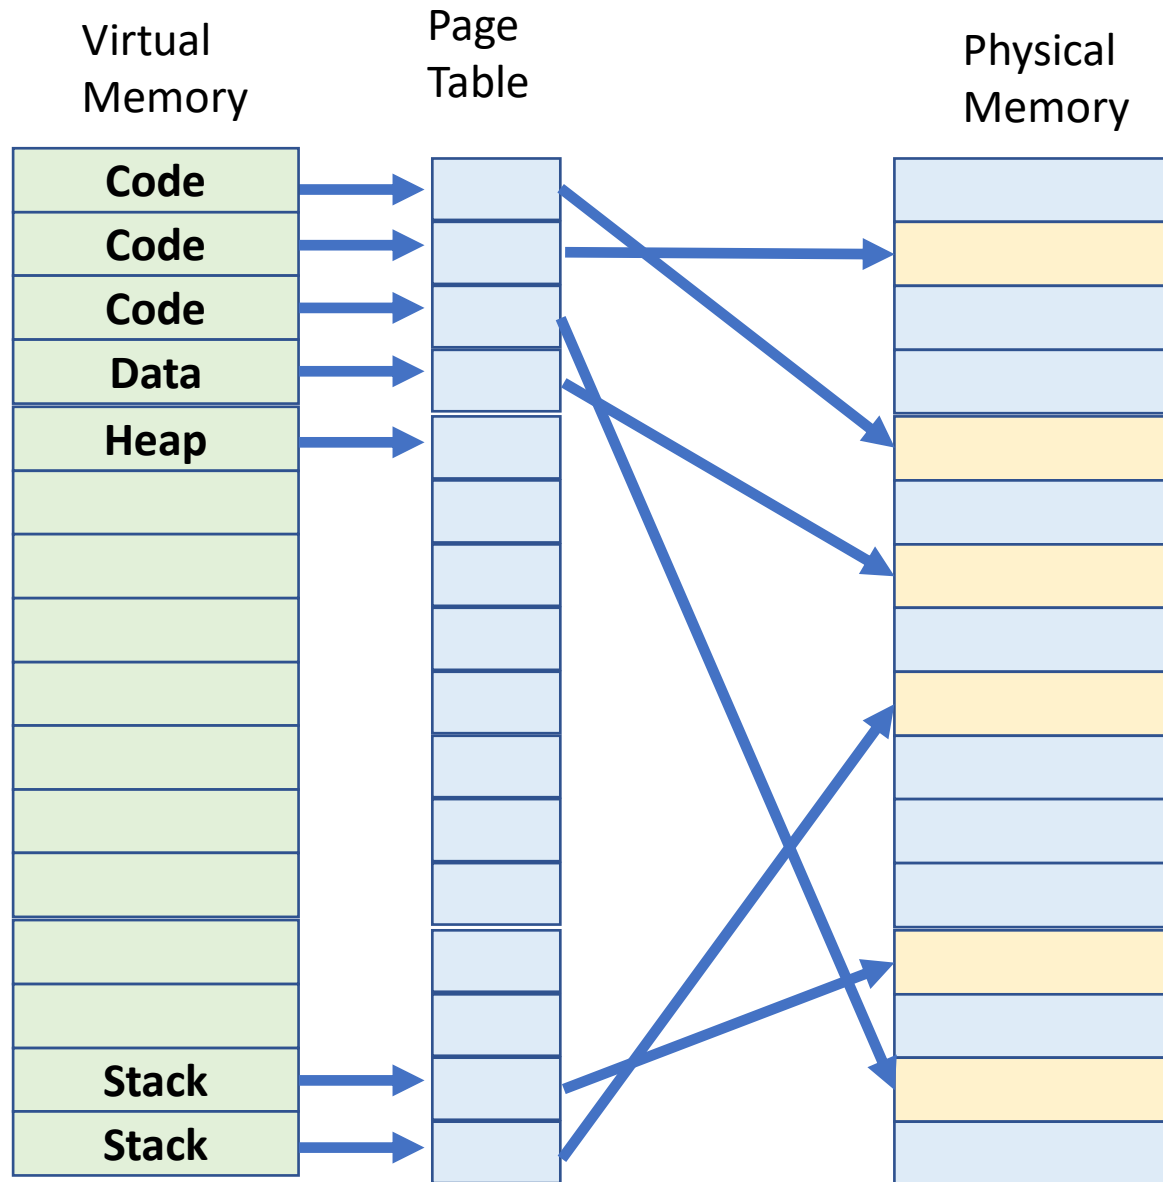
Segmentation

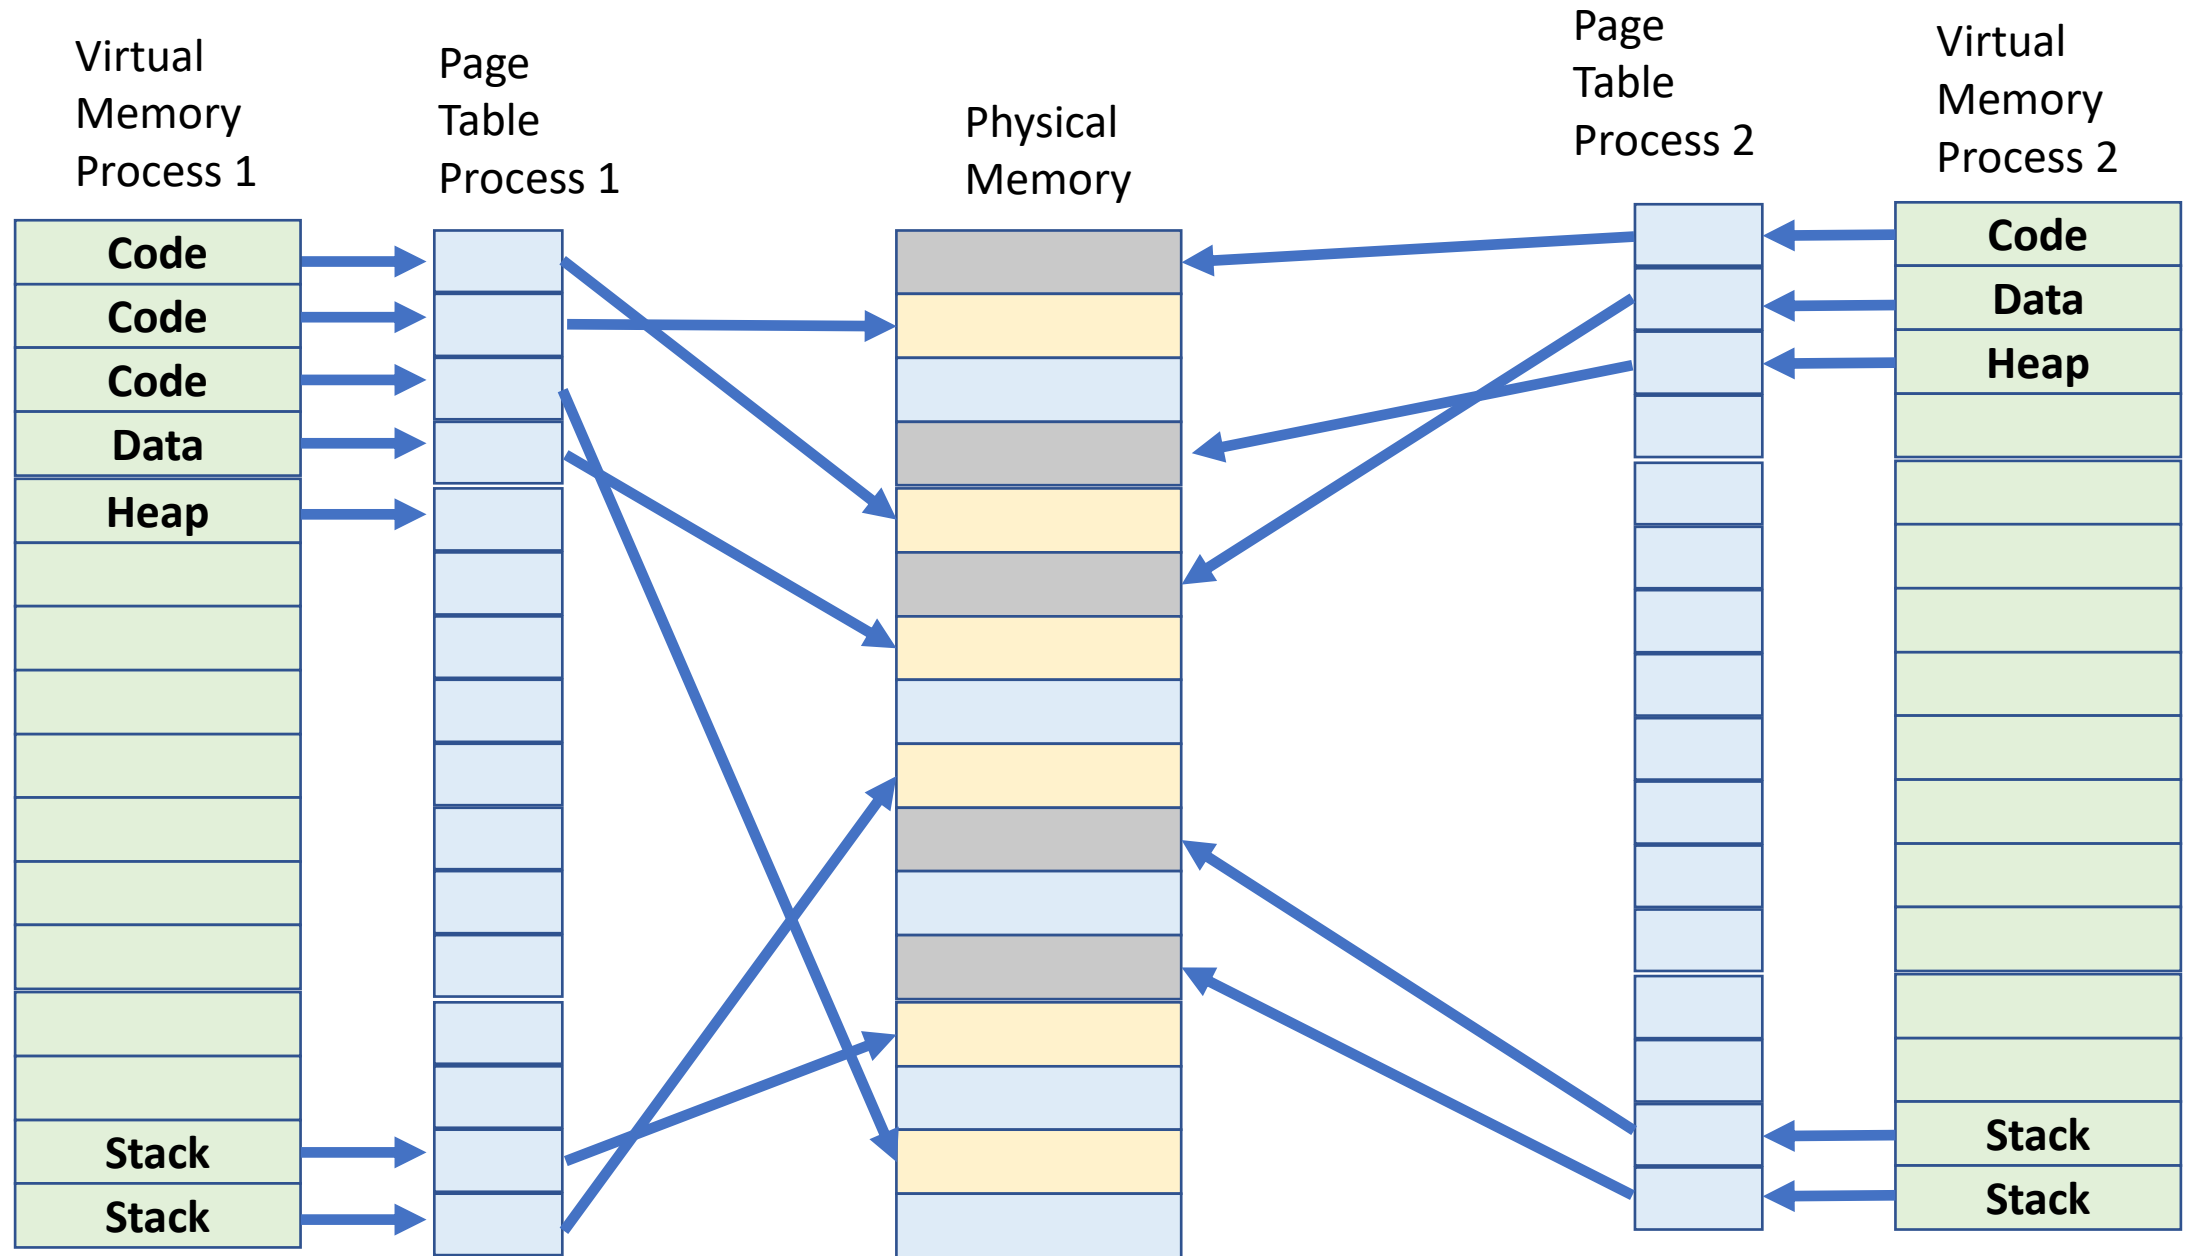


Segmentation

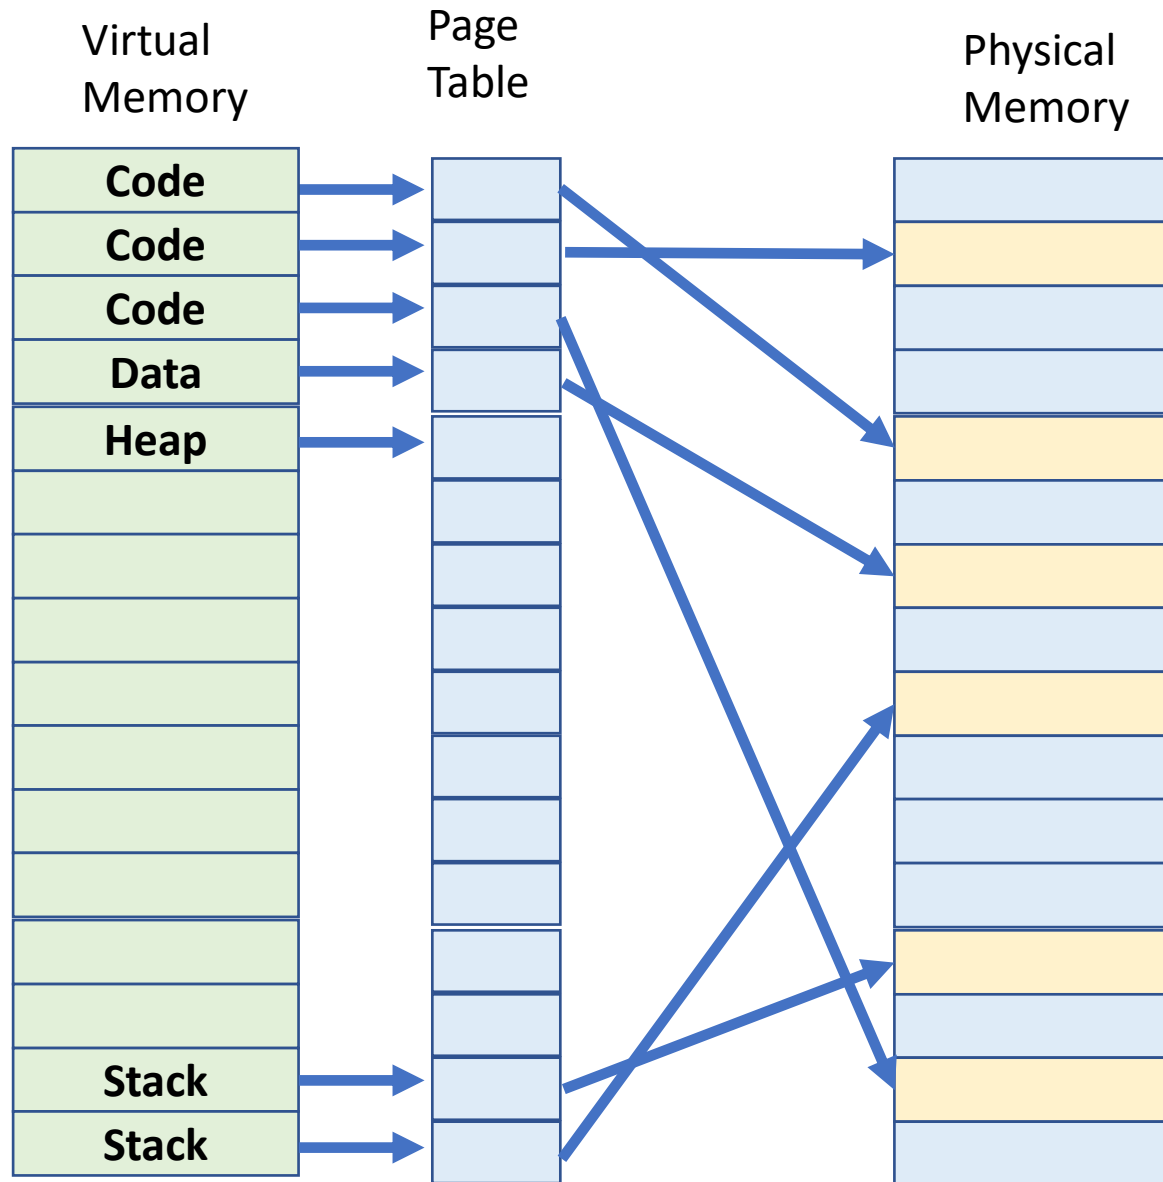
- Contiguous physical address space for each process
- Base and Bound Registers
- $\text{Physical Address} = \text{Virtual Address} + \text{Base}$
- (Must be less than bound)
- Relocatable
- External Fragmentation

Paging





Paging



- Memory divided into fixed size blocks called pages
- Example
 - page = 4096 bytes
 - 32 bit machine
 - address 0x12345678
 - page = 12345
 - offset = 678
- Each Process gets it's own page table
- Page table entries = 4 bytes
- Size of individual page table
- 2^{20} entries * 4 bytes = 4 MB
- 100 processes = 400 MB!!!

CS 354

Machine Organization and Programming

Michael Doescher
Spring 2021

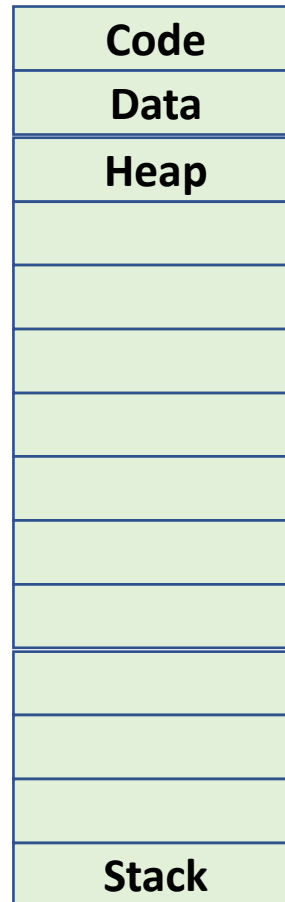
Intro to Operating Systems II

- CPU Virtualization
- Memory Virtualization
- Concurrency / Threads

Multithreading

Run multiple copies of the same code at one time

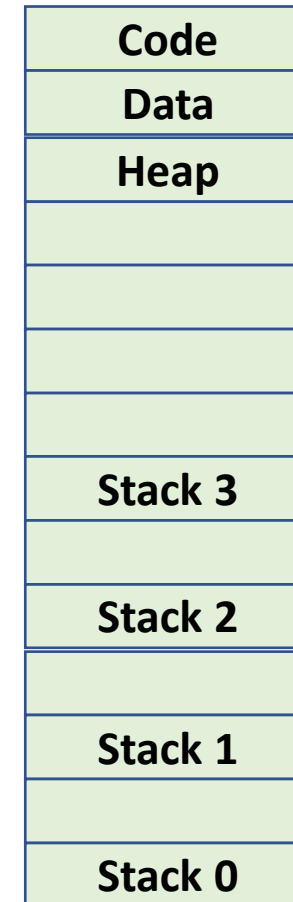
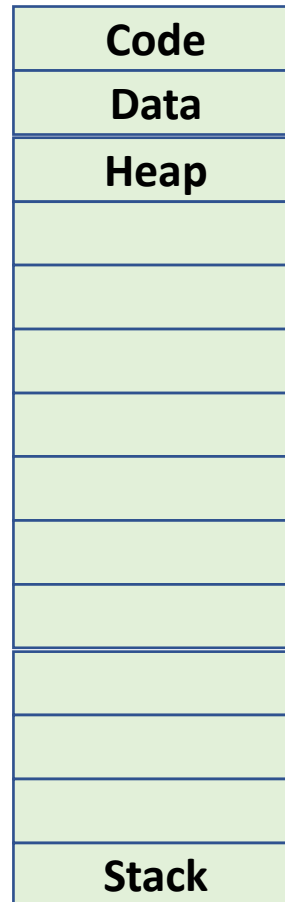
Example: Parallel processing to perform same operation on different segments of data set



Multithreading

Run multiple copies of the same code at one time

Example: Parallel processing to perform same operation on different segments of data set



Multithreading

Run multiple copies of the same code at one time

Example: Parallel processing to perform same operation on different segments of data set

Each thread gets a thread control block – and is managed (scheduled) by the OS

May execute in any order

Writing to the same memory can be a problem

“Race condition”

Atomic Instructions

Synchronization Primitives to force one thread to wait for another to finish writing

