

HASH INDEXES

CS 564- Fall 2021

ACKs: Dan Suciu, Jignesh Patel, AnHai Doan

WHAT IS THIS LECTURE ABOUT?

Hash indexes

- Static Hashing
 - I/O cost
 - issues with static hashing
- Extendible Hashing
 - insertion
 - deletion

HASH INDEXES

- efficient for equality search
- not appropriate for range search
- Types of hash indexes:
 - static hashing
 - extendible (dynamic) hashing

STATIC HASHING

- A hash index is a collection of *buckets*
 - bucket = primary page + overflow pages
 - each bucket contains one or more index entries
- To find the bucket for each record, apply a hash function *h*
- *h* maps a search key value to one of the buckets

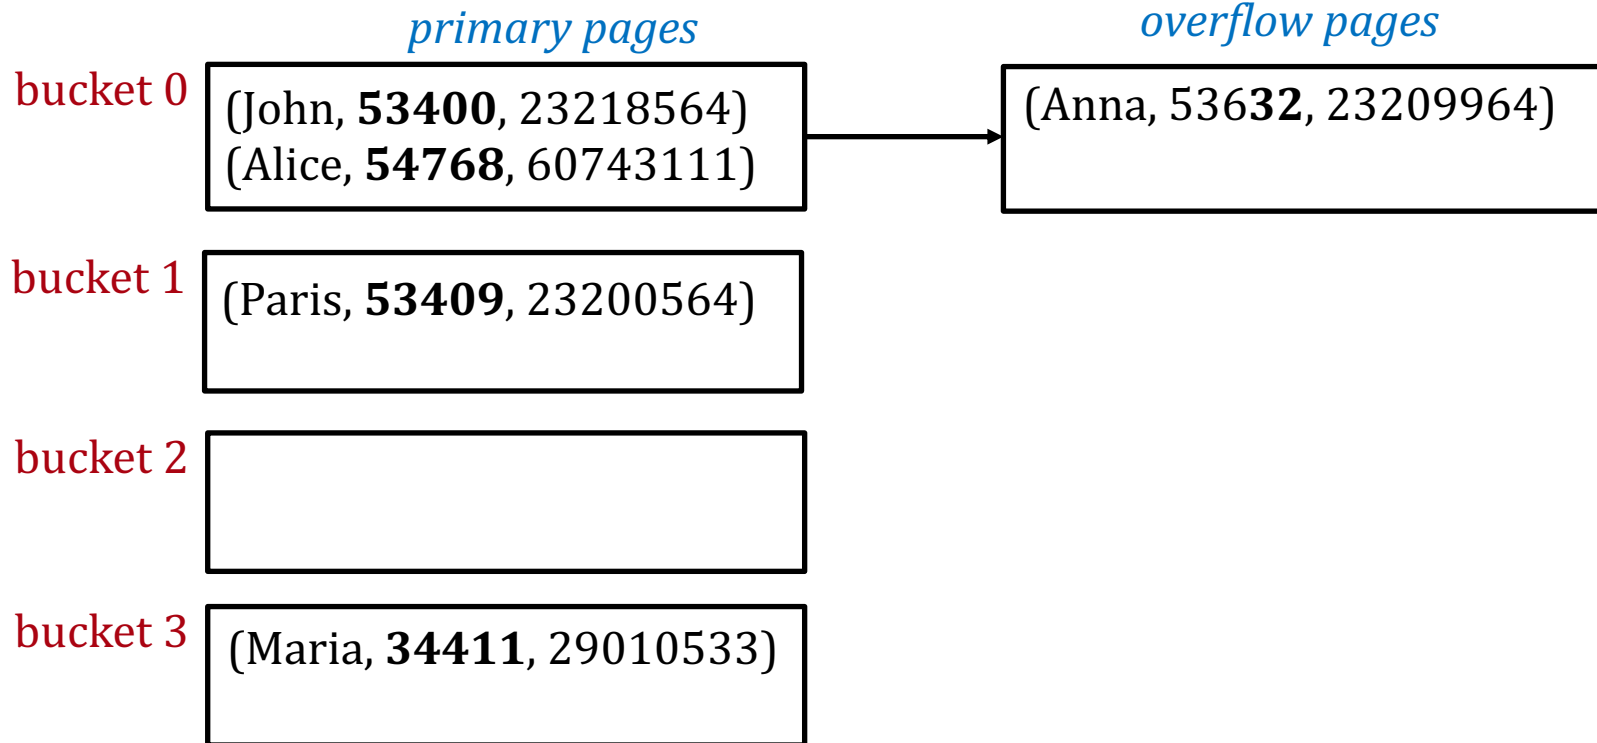
STATIC HASHING: EXAMPLE

Person(name, zipcode, phone)

- *search key*: {zipcode}
- *hash function* ***h***: zipcode mod 4

- 4 buckets

- each bucket holds 2 index entries



OPERATIONS ON HASH INDEXES

Equality search (*search-key = value*)

- apply the hash function on the search key to locate the appropriate bucket
- search through the primary page + overflow pages to find the record(s)

$$\text{I/O cost} = 1 + \text{\#overflow pages}$$

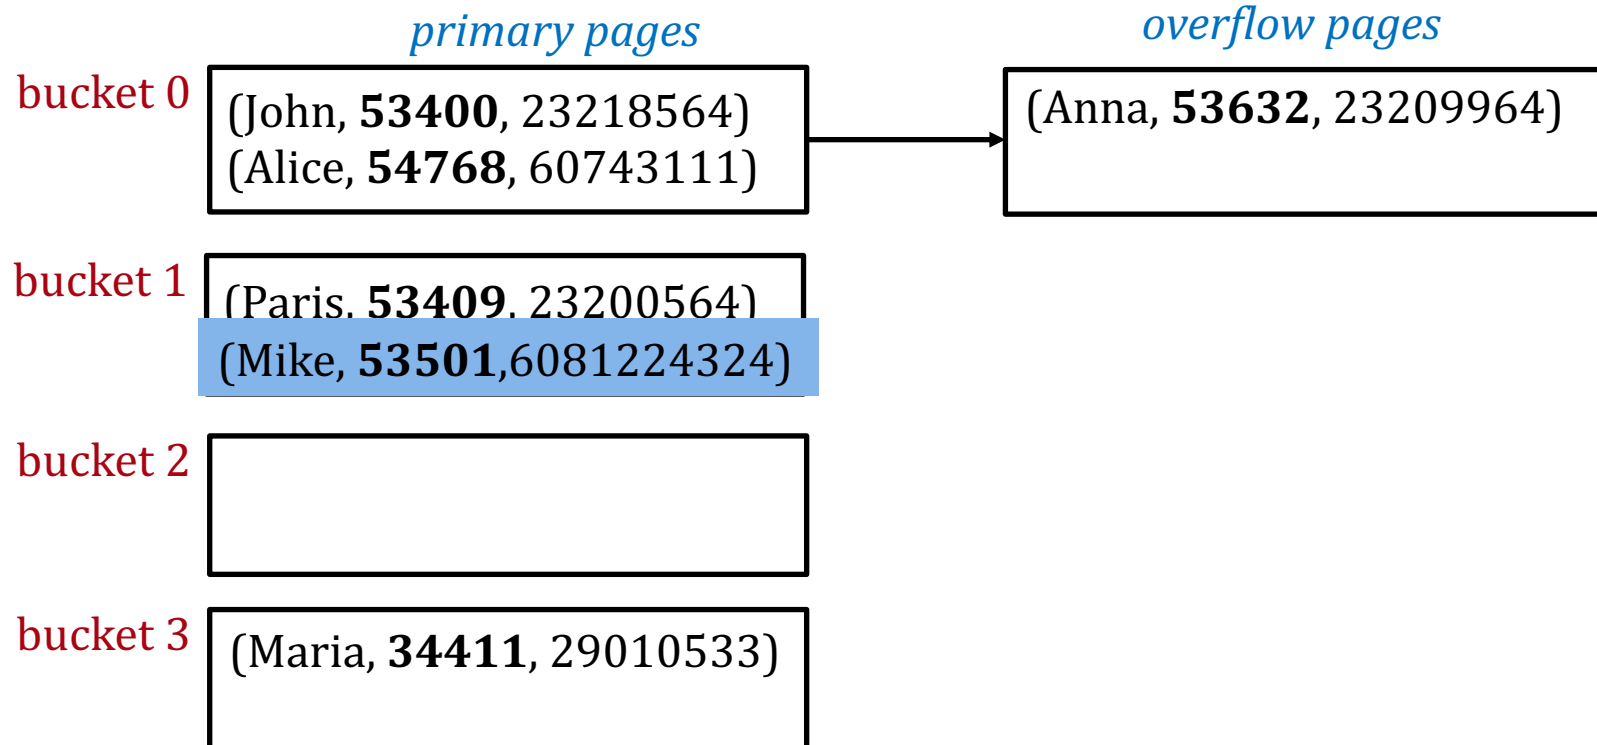
The I/O cost increases if we choose alternative #2

OPERATIONS ON HASH INDEXES

- **Deletion**
 - find the appropriate bucket, delete the record
- **Insertion**
 - find the appropriate bucket, insert the record
 - if there is no space, create a new overflow page

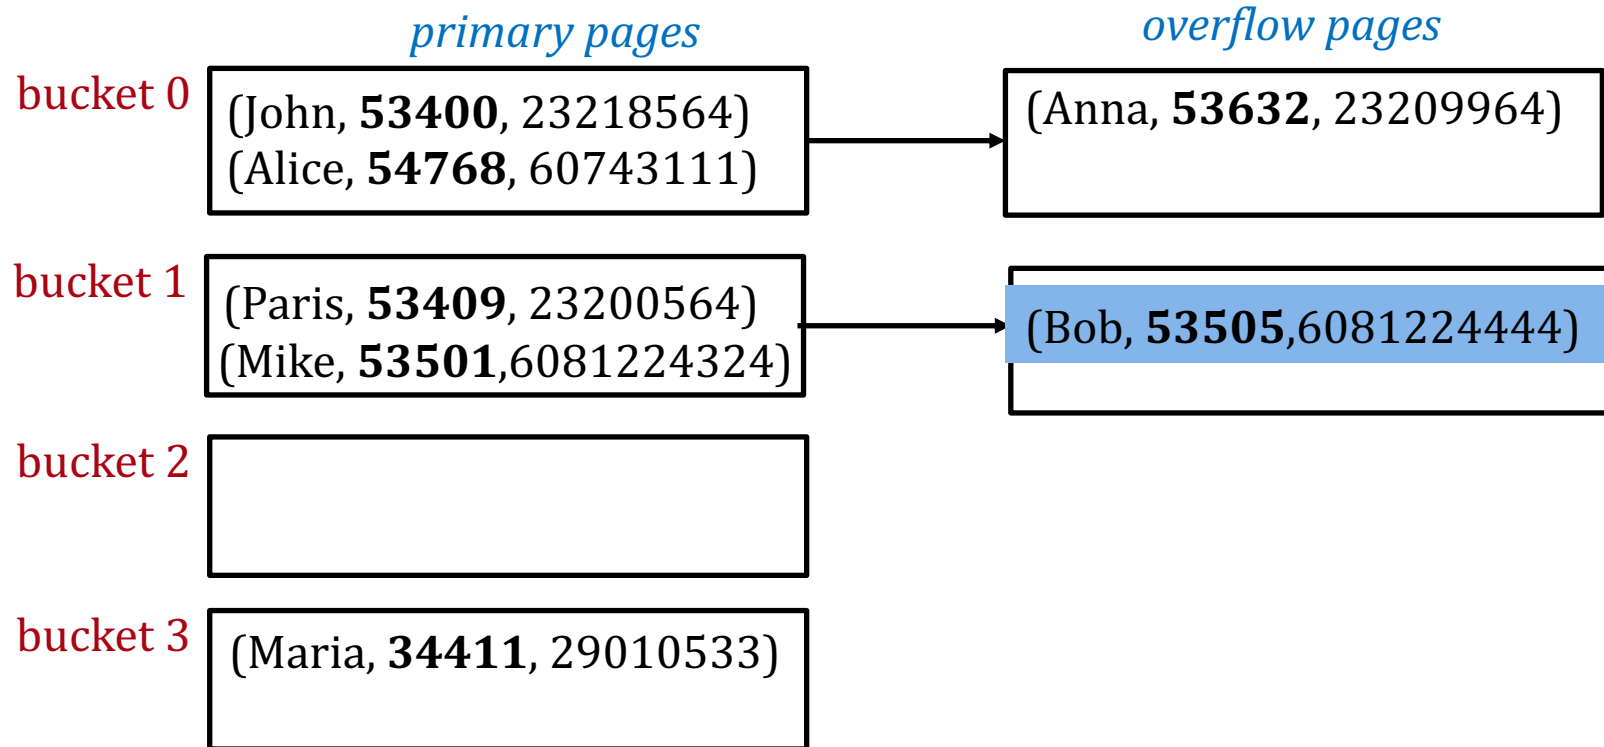
STATIC HASHING: INSERT

insert: (Mike, 53501, 6081224324)



STATIC HASHING: INSERT

insert: (Bob, 53505, 6081224444)



HASH FUNCTIONS

- An *ideal* hash function is **uniform**: each bucket is assigned the same number of key values
- A *bad* hash function maps all search key values to the same bucket
- Hash functions in practice:
 - should be fast to compute
 - should have a low collision rate

PROBLEMS OF STATIC HASHING

- In static hashing, there is a **fixed** number of buckets in the index
- Issues with this:
 - if the database grows, the number of buckets will be too small: long overflow chains degrade performance
 - if the database shrinks, space is wasted because of empty buckets
 - reorganizing the index is expensive and can block query execution

PROBLEMS OF STATIC HASHING

- Even with a good hash function, long overflow chains can still occur when we have **skew**
- Skew occurs when many records have the same search key value, so they will always end up in the same bucket

EXTENDIBLE HASHING

EXTENDIBLE HASHING

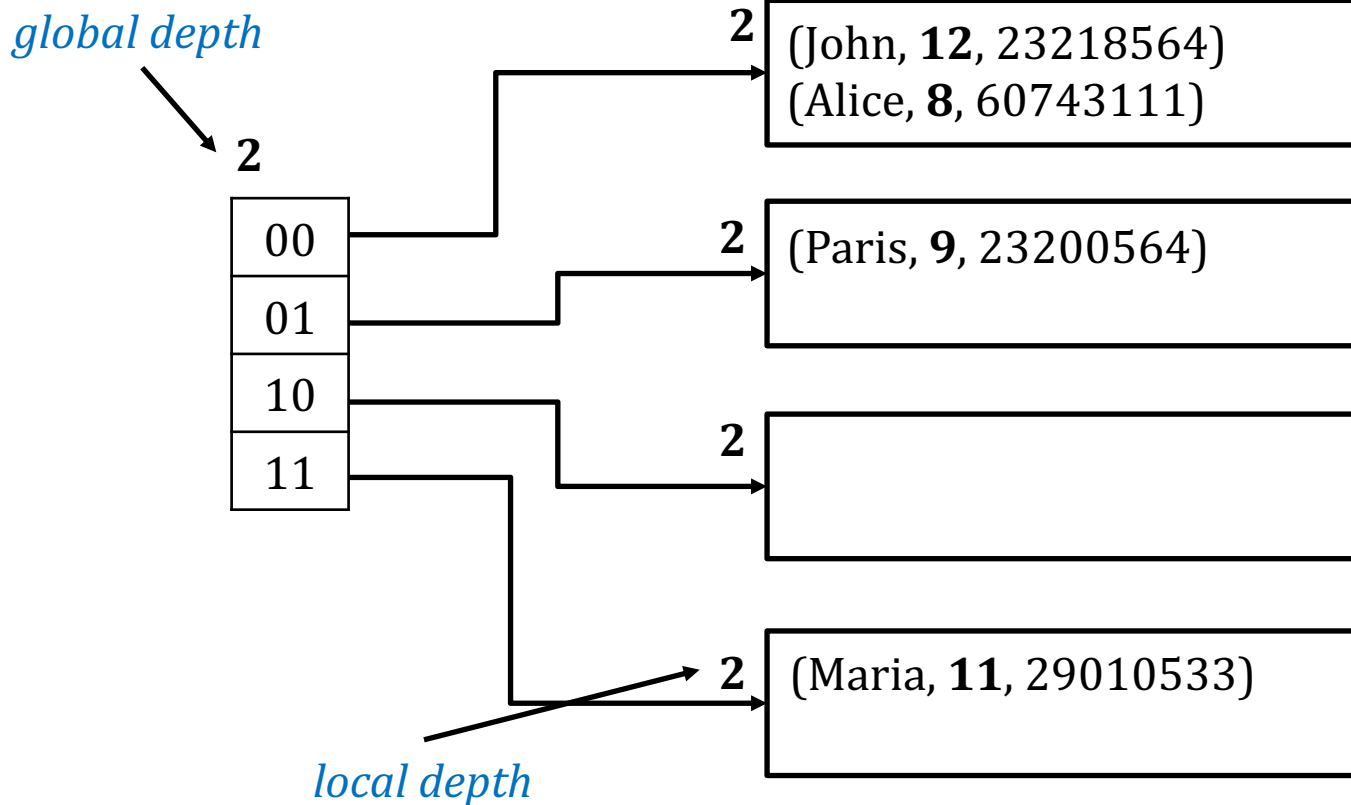
Extendible hashing is a type of *dynamic* hashing

- It keeps a directory of pointers to buckets
- On overflow, it reorganizes the index by **doubling the directory** (and not the number of buckets)

EXTENDIBLE HASHING

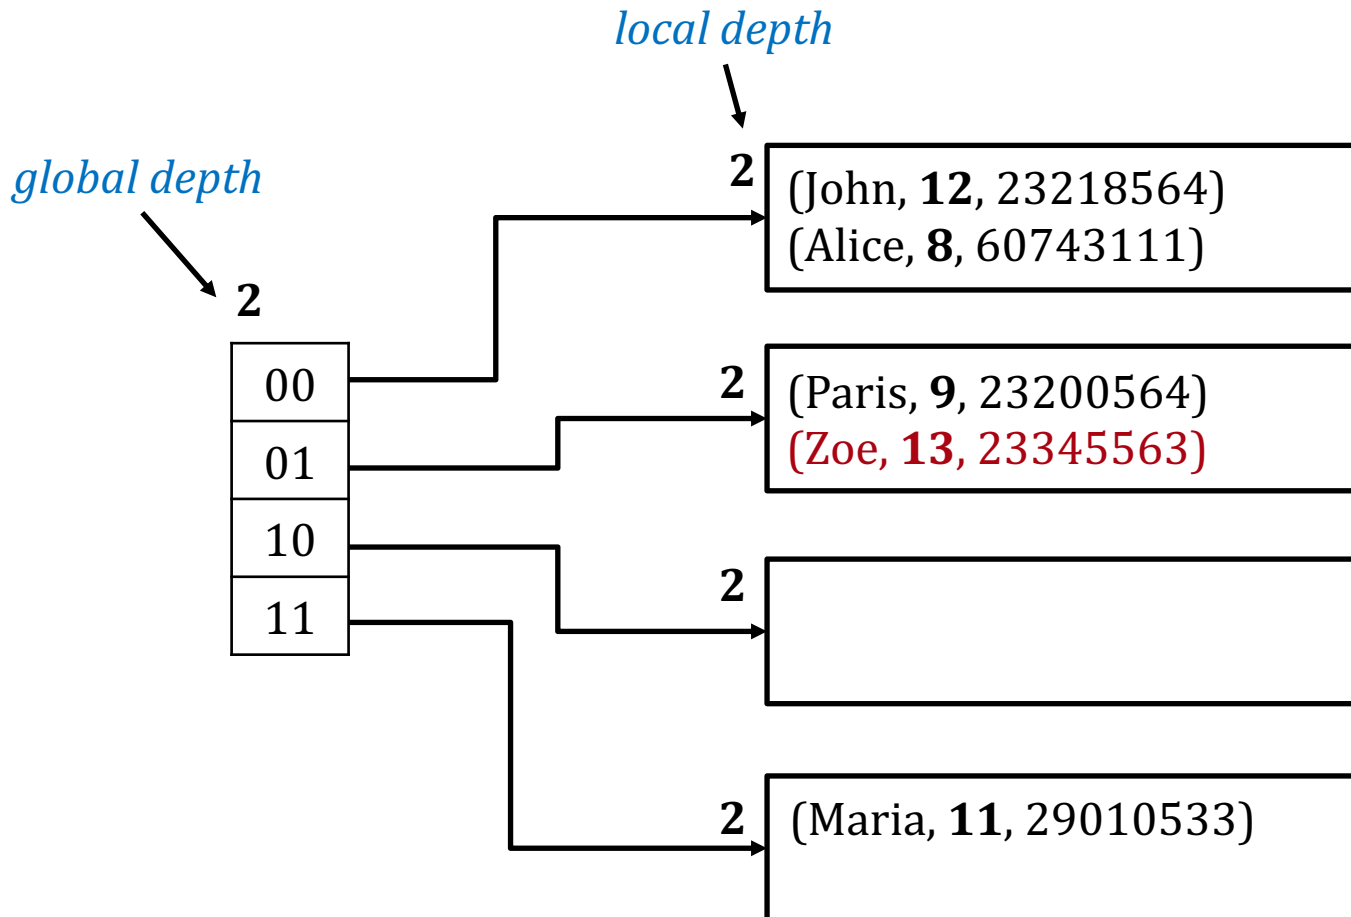
To search, use the last (*global depth*) digits of the **binary** form of the search key value

12 = 1100



EXTENDIBLE HASHING: INSERT

If there is space in the bucket, simply add the record




EXTENDIBLE HASHING: INSERT

If the bucket is full, split the bucket and redistribute the entries

global depth increases by 1

3




local depth increases for the split bucket!

local depth remains the same for the other buckets

EXAMPLE

each page can hold at most two records




INVARIANT: global depth \geq local depth

EXAMPLE

- The catalog doubles in size
- Global depth becomes 2

insert: (22,...)



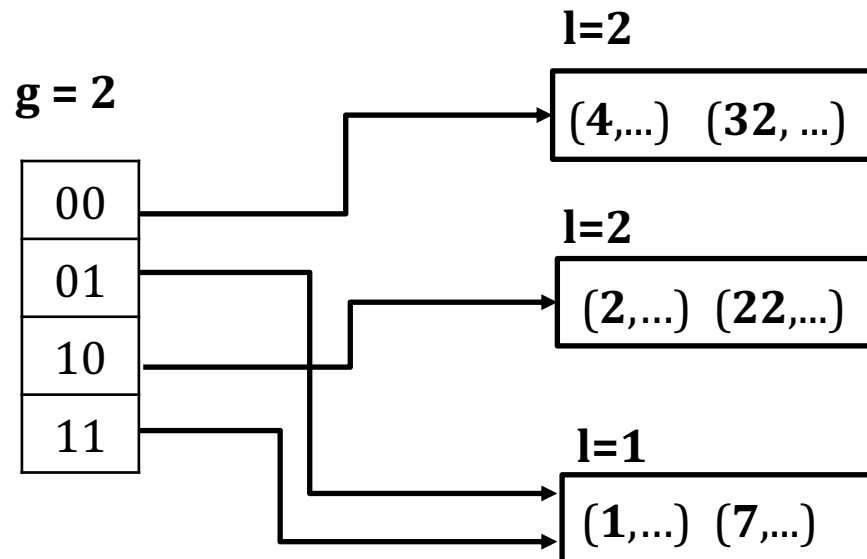
This bucket is split into two buckets with local depth 2

This bucket remains the same

EXAMPLE

There is space in the bucket
so nothing changes!

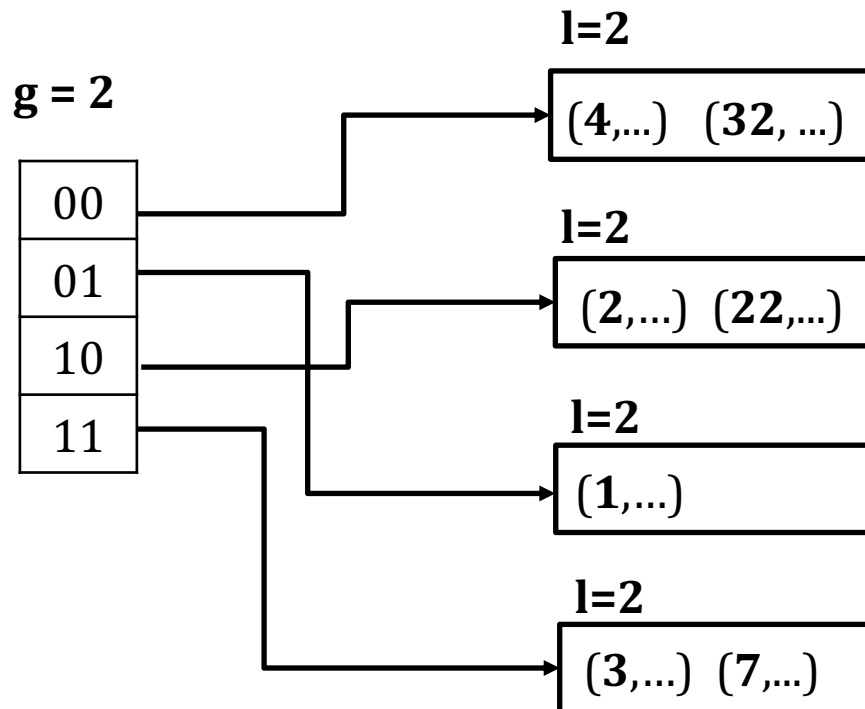
insert: (32,...)



EXAMPLE

Since local depth is smaller than global,
no need to change the directory size!

insert: (3,...)



The bucket is split into two
buckets with local depth 2

EXTENDIBLE HASHING: DELETE

- Locate the bucket of the record and remove it
- If the bucket becomes empty, it can be removed (and update the directory)
- Two buckets can also be coalesced together if the sum of the entries fit in a single bucket
- Decreasing the size of the directory can also be done, but it is expensive

MORE ON EXTENDIBLE HASHING

- How many I/Os for equality search?
 - One if directory fits in memory, else two
- The directory grows in spurts, and, if the distribution of hash values is skewed, the directory can grow very large
- We may need overflow pages when multiple entries have the same hash value!