

# RELATIONAL OPERATORS: OTHER

---

*CS 564- Fall 2021*

---

*ACKs: Jeff Naughton, Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT?

---

We will discuss algorithms for

- Projection
- Set operations
- Aggregation

---

# PROJECTION

---

---

# PROJECT OPERATOR

---

**Simple case:** **SELECT R.A, R.D**

- scan the file and for each tuple output R.A, R.D

**Hard case:** **SELECT DISTINCT R.A, R.D**

- project out the attributes
- eliminate *duplicate tuples* (this is the difficult part!)

---

# PROJECT: SORT-BASED

---

## **Naïve algorithm:**

1. scan the relation and project out the attributes
2. sort the resulting set of tuples using all remaining attributes
3. scan the sorted set by comparing only adjacent tuples and discard duplicates

---

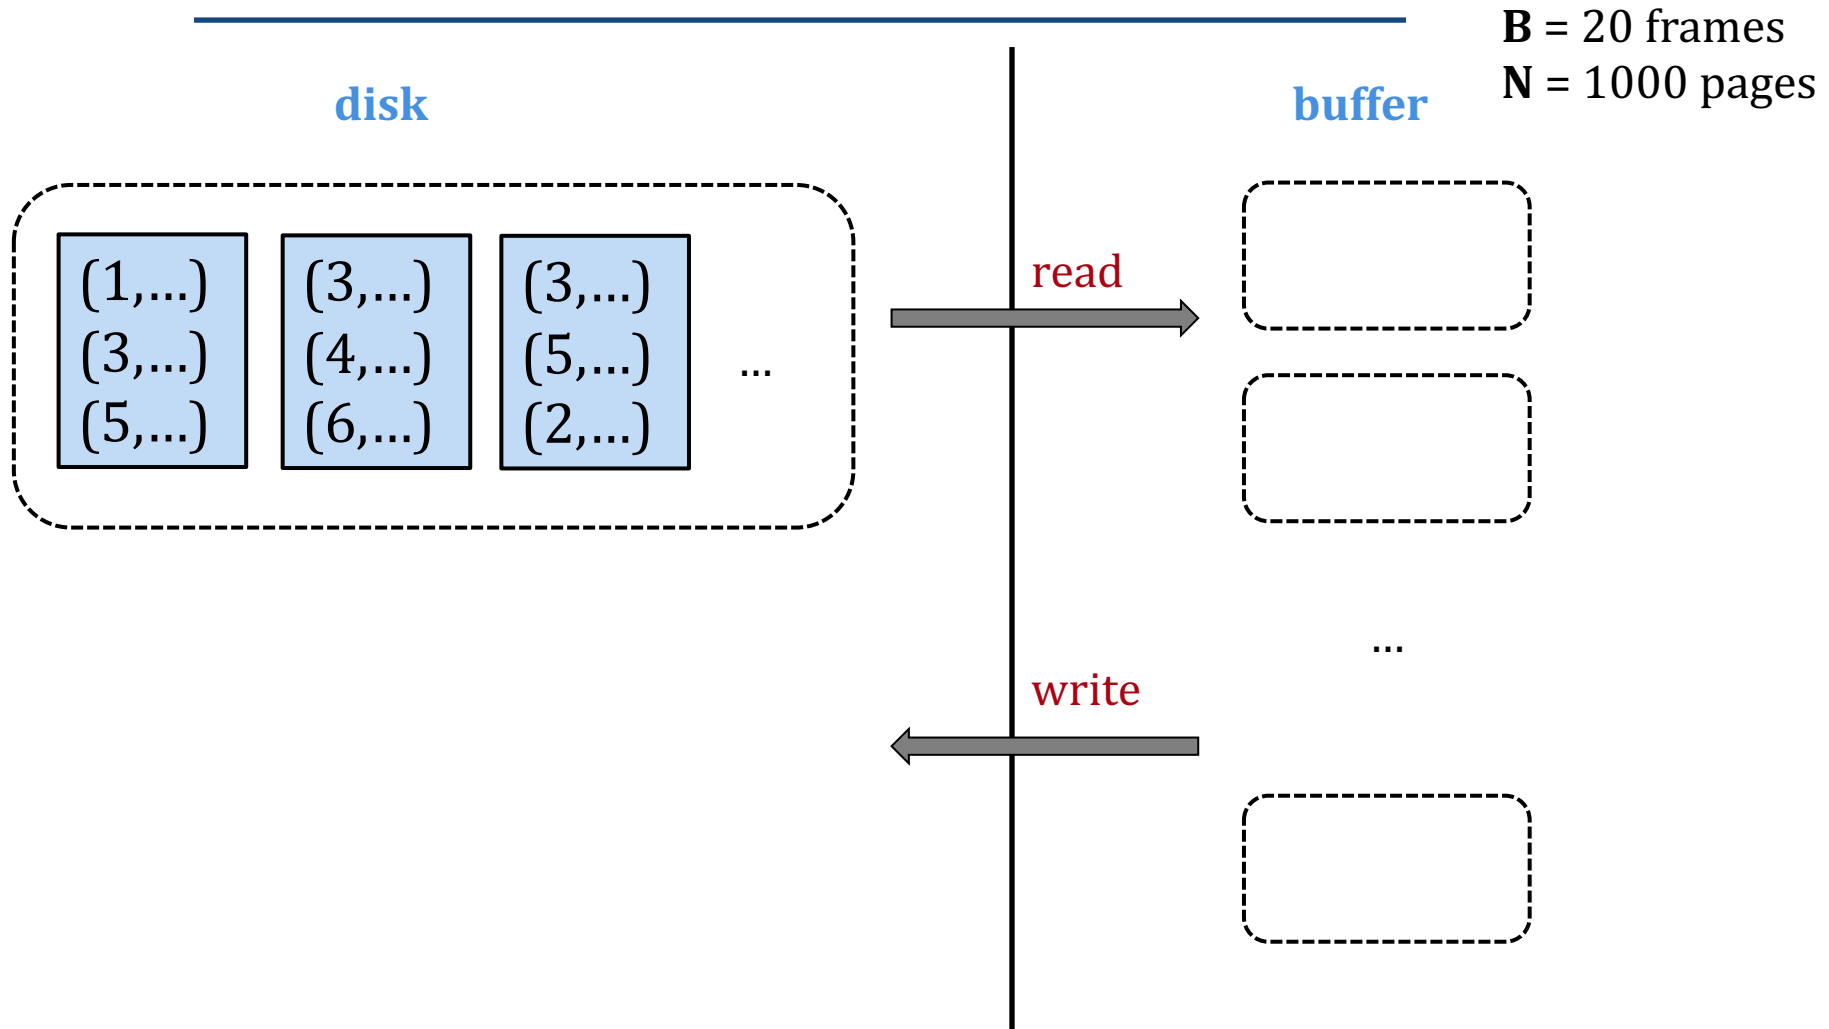
# RUNNING EXAMPLE

---

$R(A, B, C, D, E)$

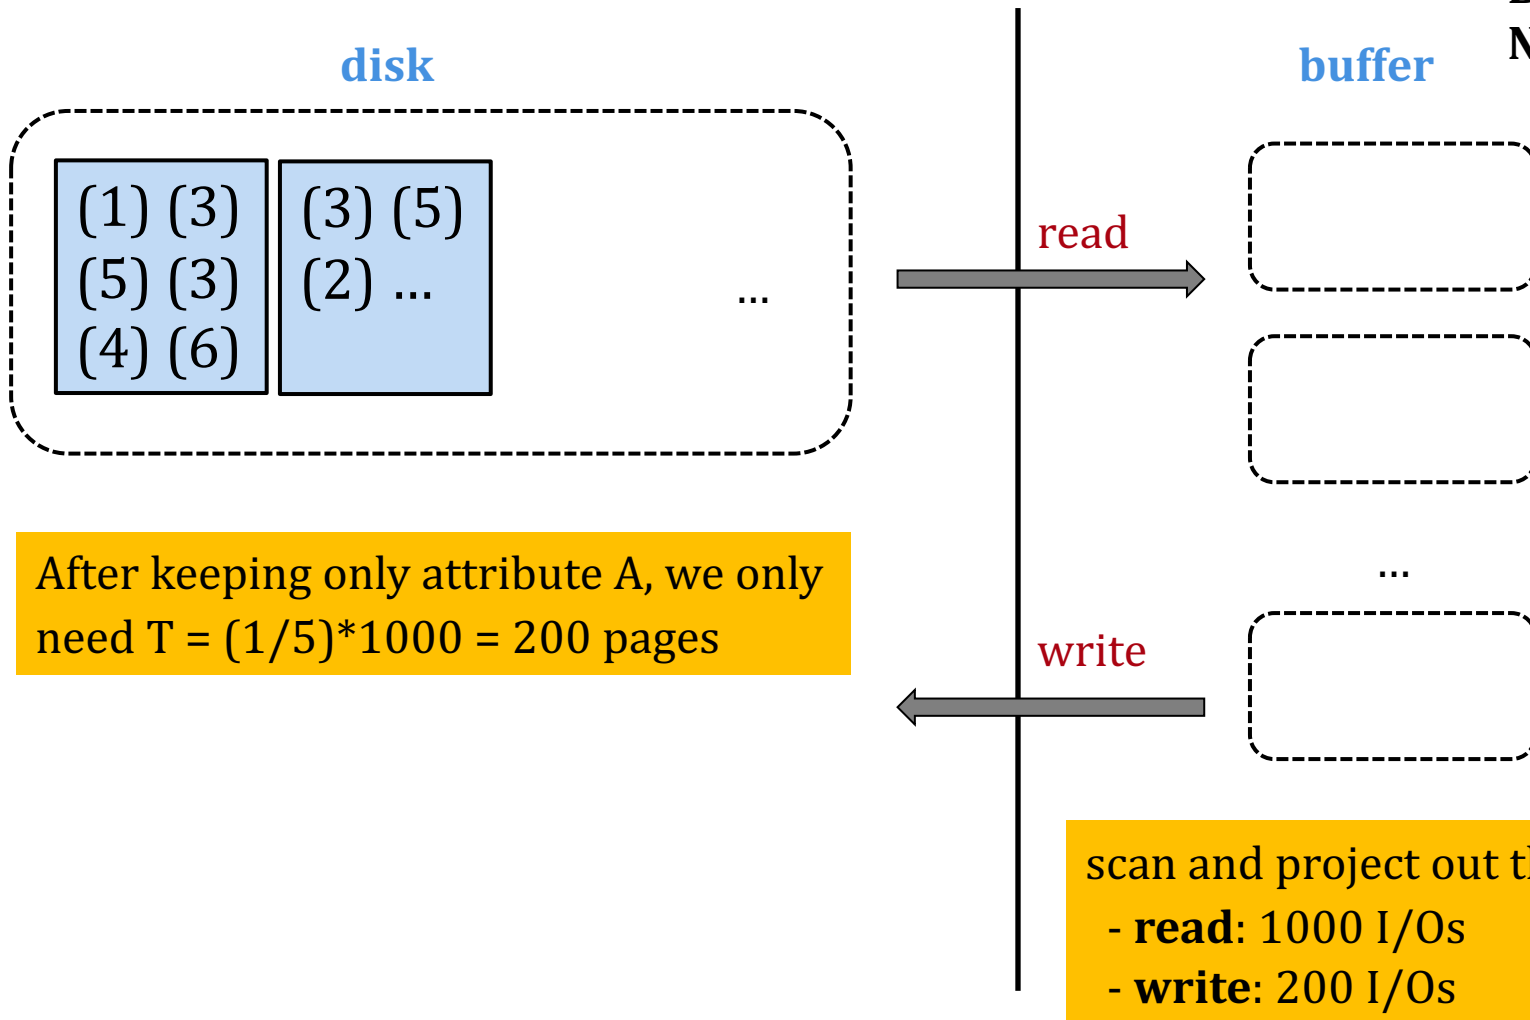
- $N = 1000$  pages
- $B = 20$  buffer pages
- Each field in the tuple has the same size
- Suppose we want to project on attribute  $A$

# RUNNING EXAMPLE



# RUNNING EXAMPLE

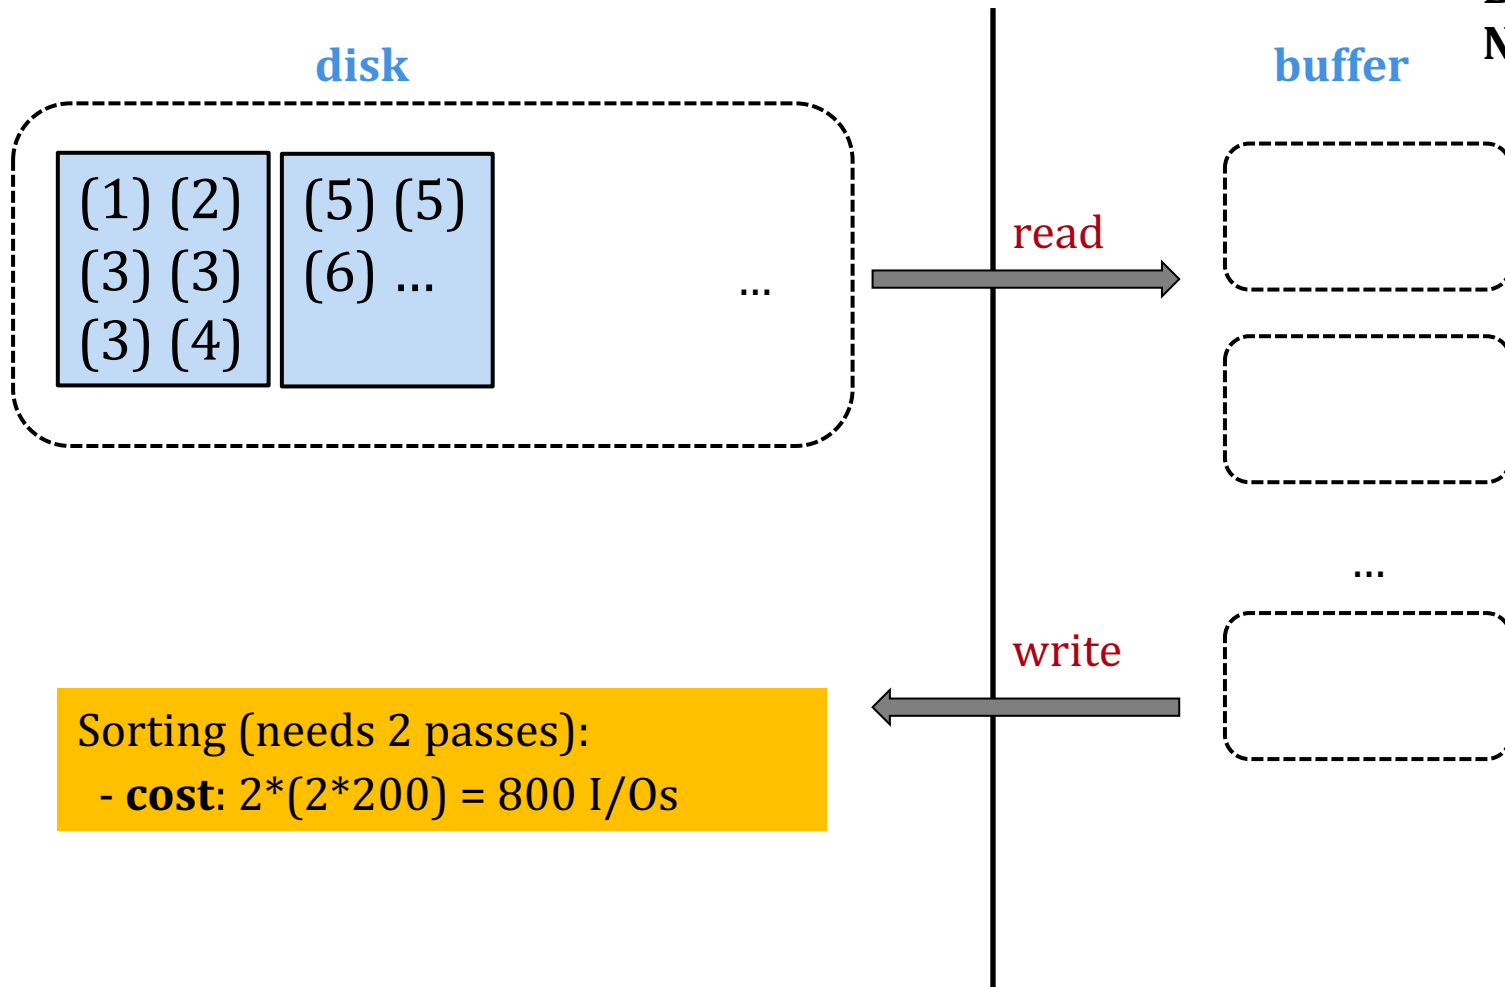
B = 20 frames  
N = 1000 pages





# RUNNING EXAMPLE

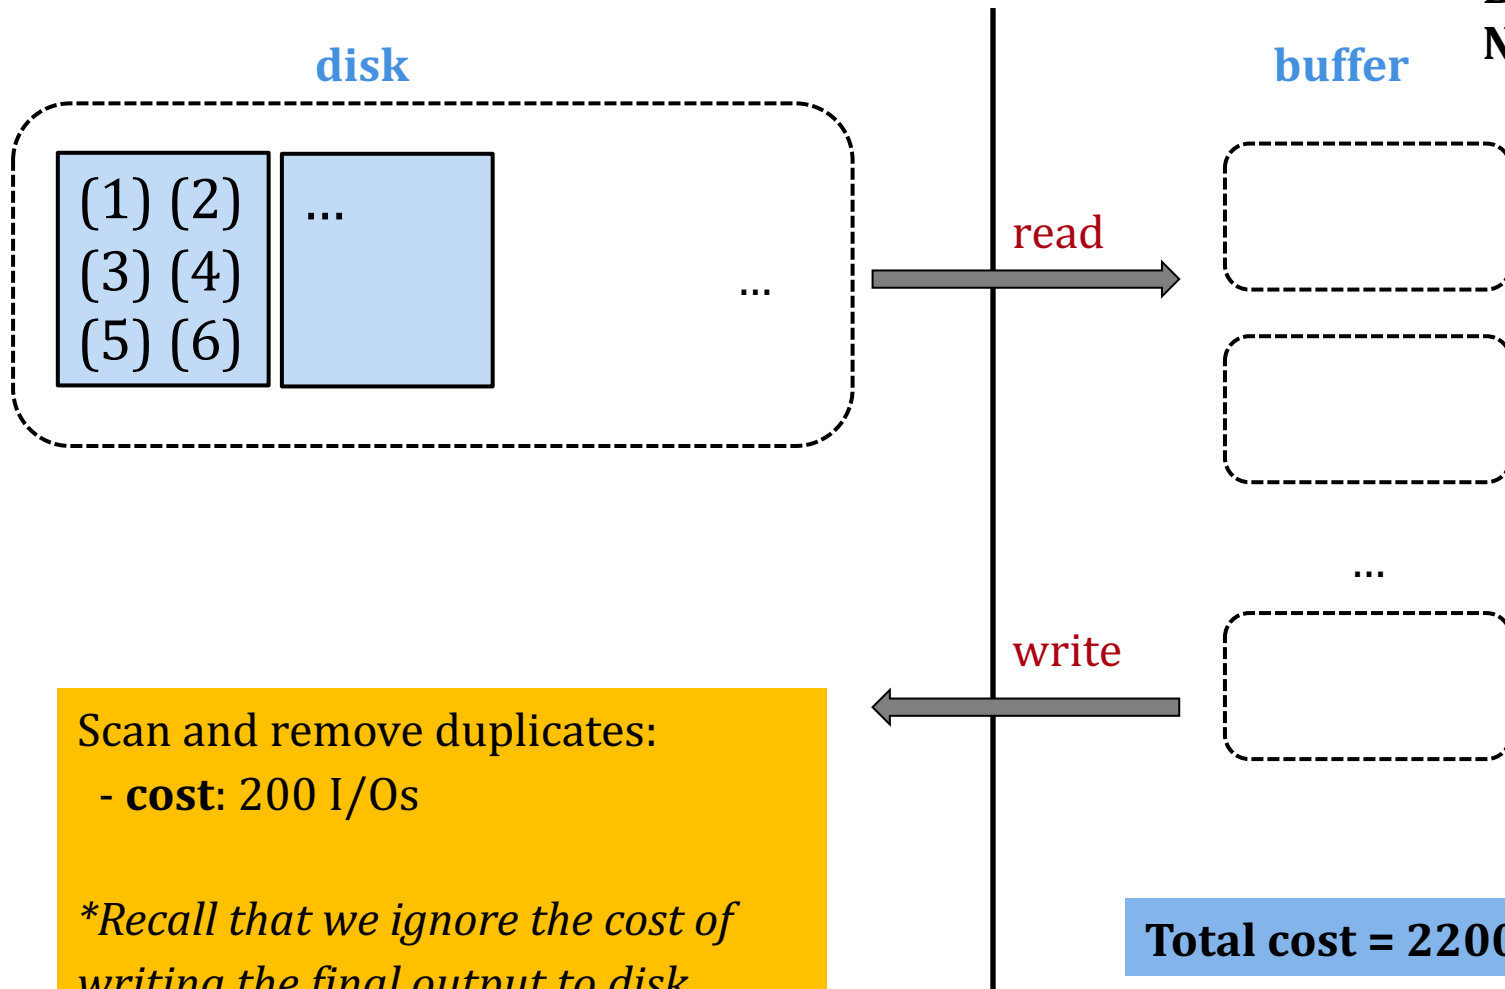
B = 20 frames  
N = 1000 pages



Sorting (needs 2 passes):  
- **cost**:  $2 \cdot (2 \cdot 200) = 800$  I/Os

# RUNNING EXAMPLE

B = 20 frames  
N = 1000 pages



Scan and remove duplicates:

- **cost**: 200 I/Os

*\*Recall that we ignore the cost of writing the final output to disk*

**Total cost = 2200 I/Os**

---

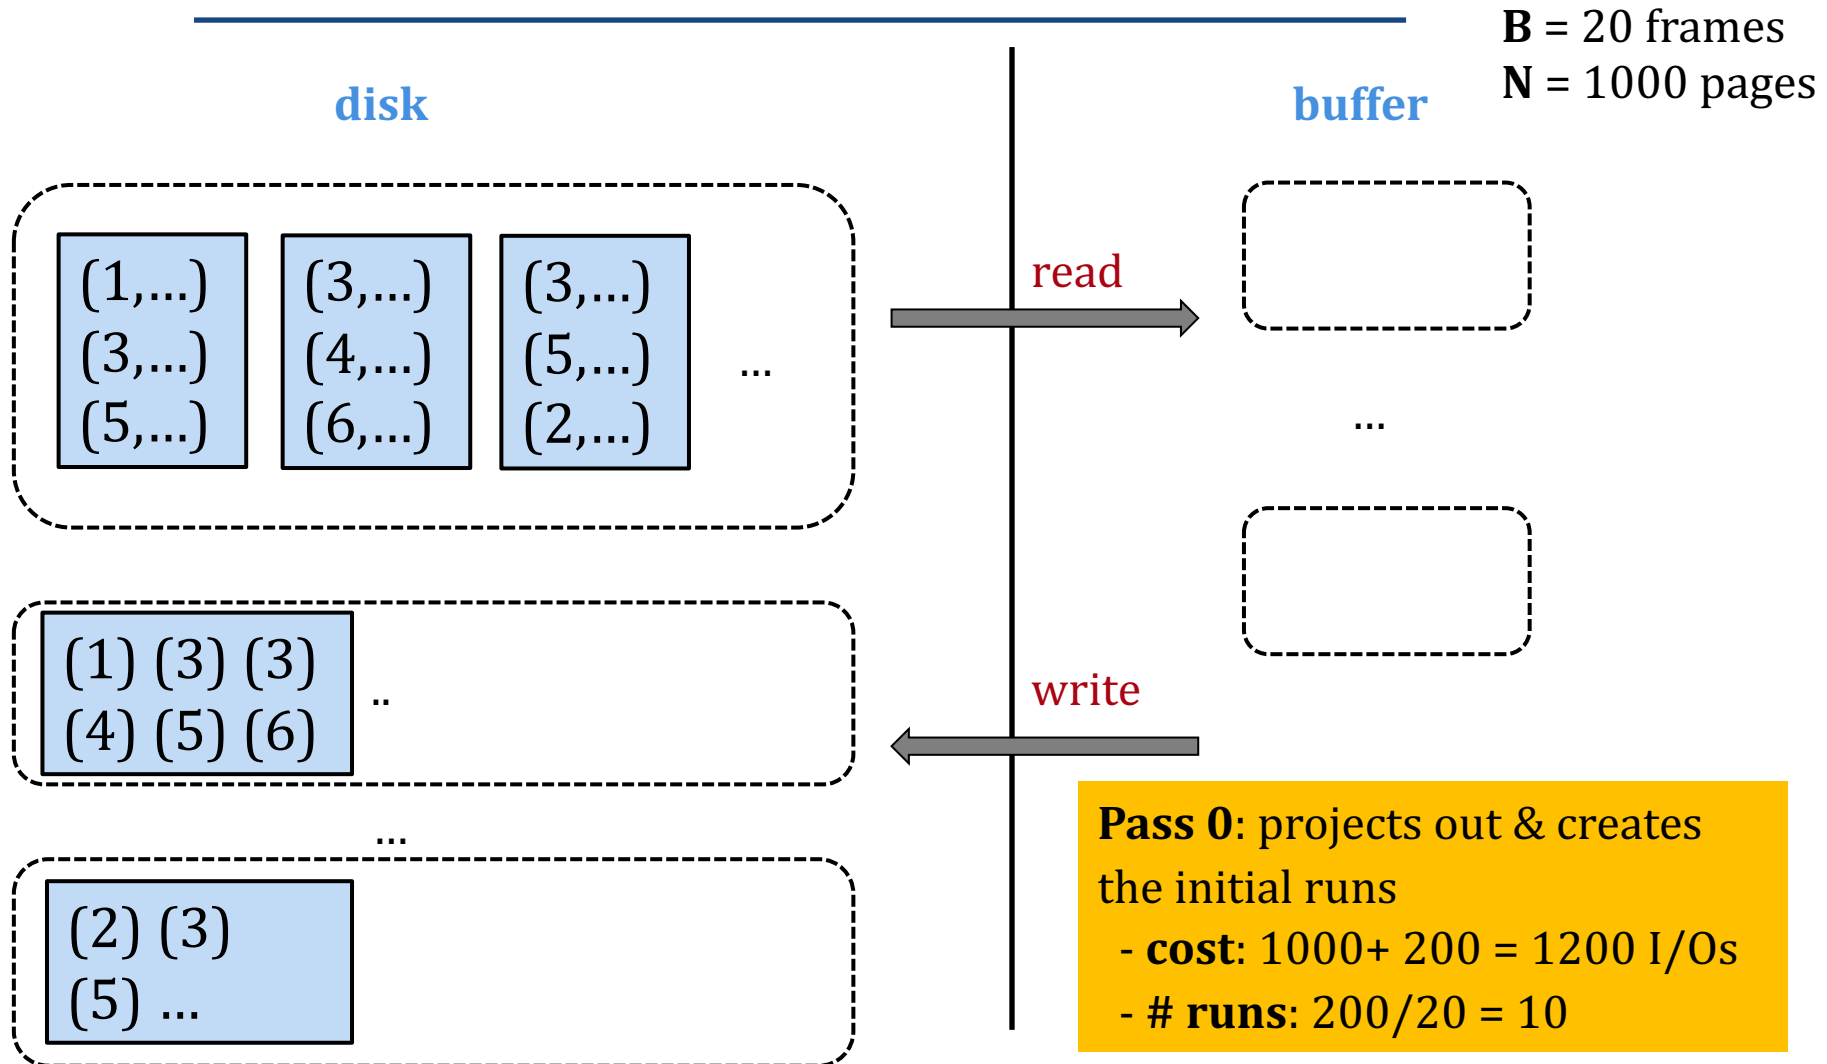
# PROJECT: SORT-BASED

---

We can improve upon the naïve algorithm by modifying the sorting algorithm:

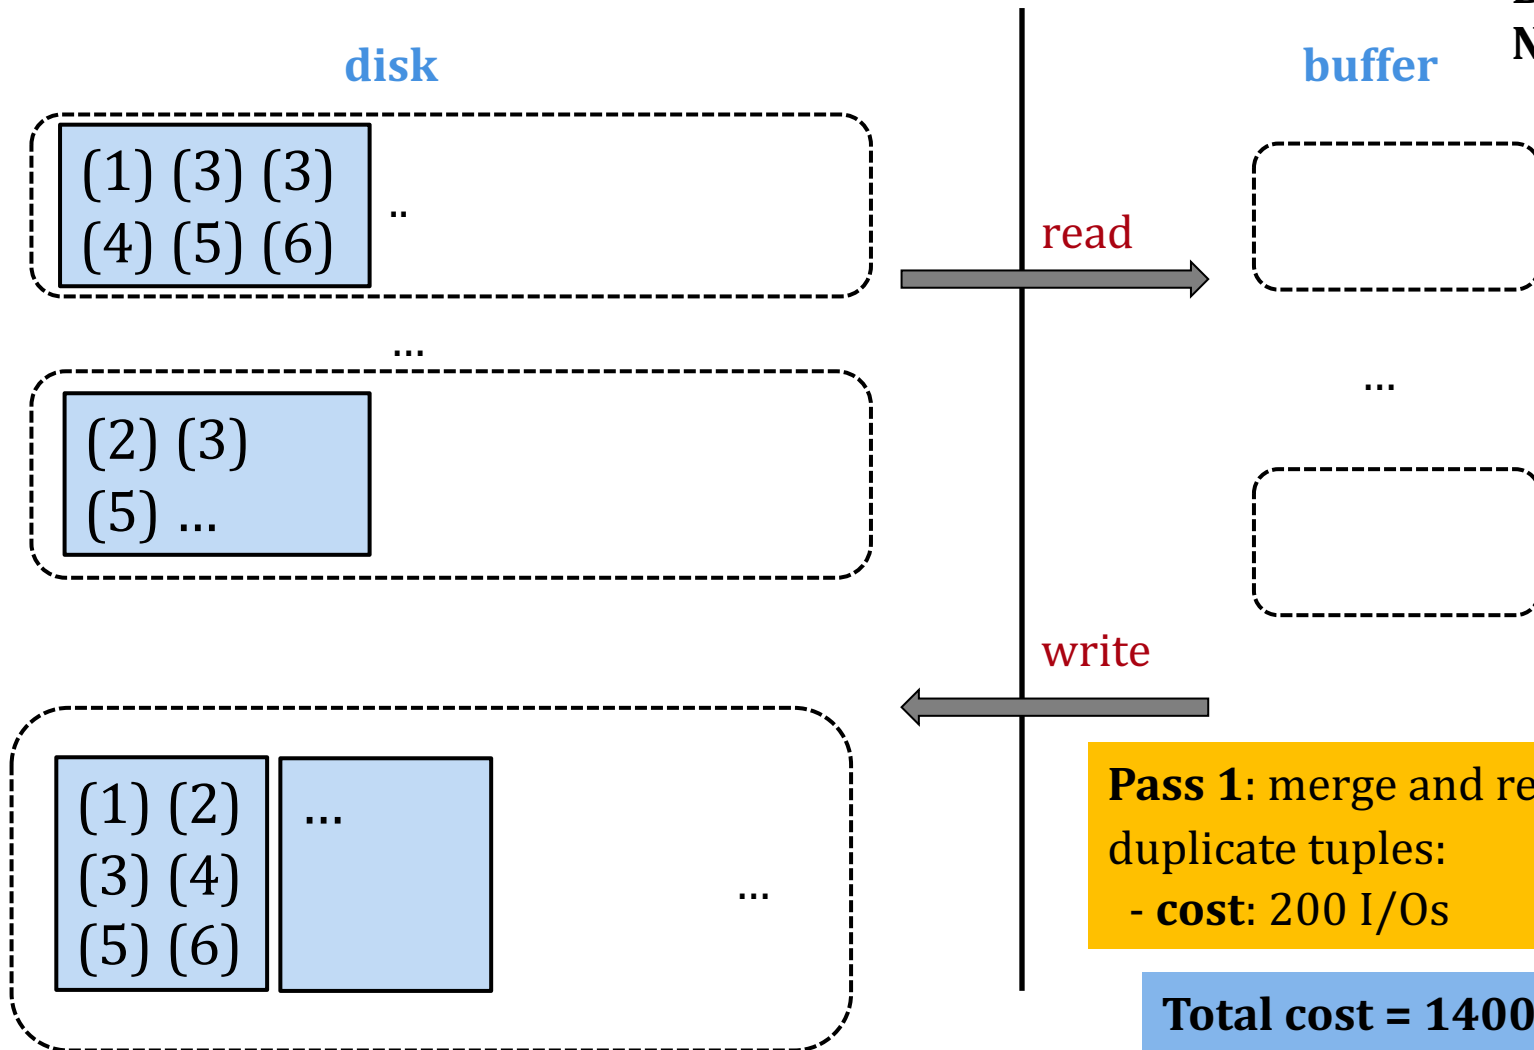
1. In Pass **0** of sorting, project out the attributes
2. In subsequent passes, eliminate the duplicates while merging the runs

# RUNNING EXAMPLE



# RUNNING EXAMPLE

B = 20 frames  
N = 1000 pages



**Pass 1:** merge and remove duplicate tuples:  
- **cost:** 200 I/Os

**Total cost = 1400 I/Os**

# PROJECT: HASH-BASED

---

2-phase algorithm:

- **partitioning**
  - project out attributes and split the input into  $B-1$  partitions using a hash function  $h$
- **duplicate elimination**
  - read each partition into memory and use an in-memory hash table (with a *different* hash function) to remove duplicates

# PROJECT: HASH-BASED

---

When does the hash table fit in memory?

- size of a partition =  $T / (B - 1)$ , where  $T$  is #pages after projection
- size of hash table =  $f \cdot T / (B - 1)$ , where  $f$  is the **fudge factor** (typically  $\sim 1.2$ )
- So, it must be  $B > f \cdot T / (B - 1)$

---

# **HASH-BASED COST ANALYSIS**

---

- $T = 200$  so the hash table fits in memory!
- partitioning cost =  $1000 + 200 = 1200$  I/Os
- duplicate elimination cost =  $200$  I/Os

**total cost = 1400 I/Os**



# COMPARISON

---

Benefits of sort-based approach

- better handling of skew
- the output result is sorted

For large values of the buffer size  $B$ , both algorithms need only 2 passes!

---

# PROJECT: INDEX-BASED

---

- Index-only scan
  - projection attributes are a subset of index attributes
  - apply projection algorithm only to index entries!
- If an *ordered index* contains all projection attributes as prefix of the search key:
  1. retrieve index data entries in order
  2. discard unwanted fields
  3. compare adjacent entries to eliminate duplicates

---

# SET OPERATIONS

---

---

# SET OPERATIONS

---

- **Intersection** is a special case of a join, so any algorithm for equijoins applies
- **Union** and **difference** are similar

---

# UNION

---

- Sort-based:
  - sort both relations (on *all attributes*)
  - merge sorted relations while eliminating duplicates
- Hash-based:
  - hash-partition  $R$  and  $S$
  - build an in-memory hash table for each partition  $R_i$
  - probe with tuples in  $S_i$ , add to the output result if not a duplicate

---

# SET DIFFERENCE

---

We want to compute  $R - S$

- Sort-based:
  - sort both relations (on *all attributes*)
  - merge sorted relations while eliminating tuples from  $R$  that exist in  $S$
- Hash-based:
  - hash-partition  $R$  and  $S$
  - build an in-memory hash table for each partition  $S_i$
  - probe with tuples in  $R_i$ , add to the output result if the tuple does not exist in the hash table

---

# AGGREGATION

---

---

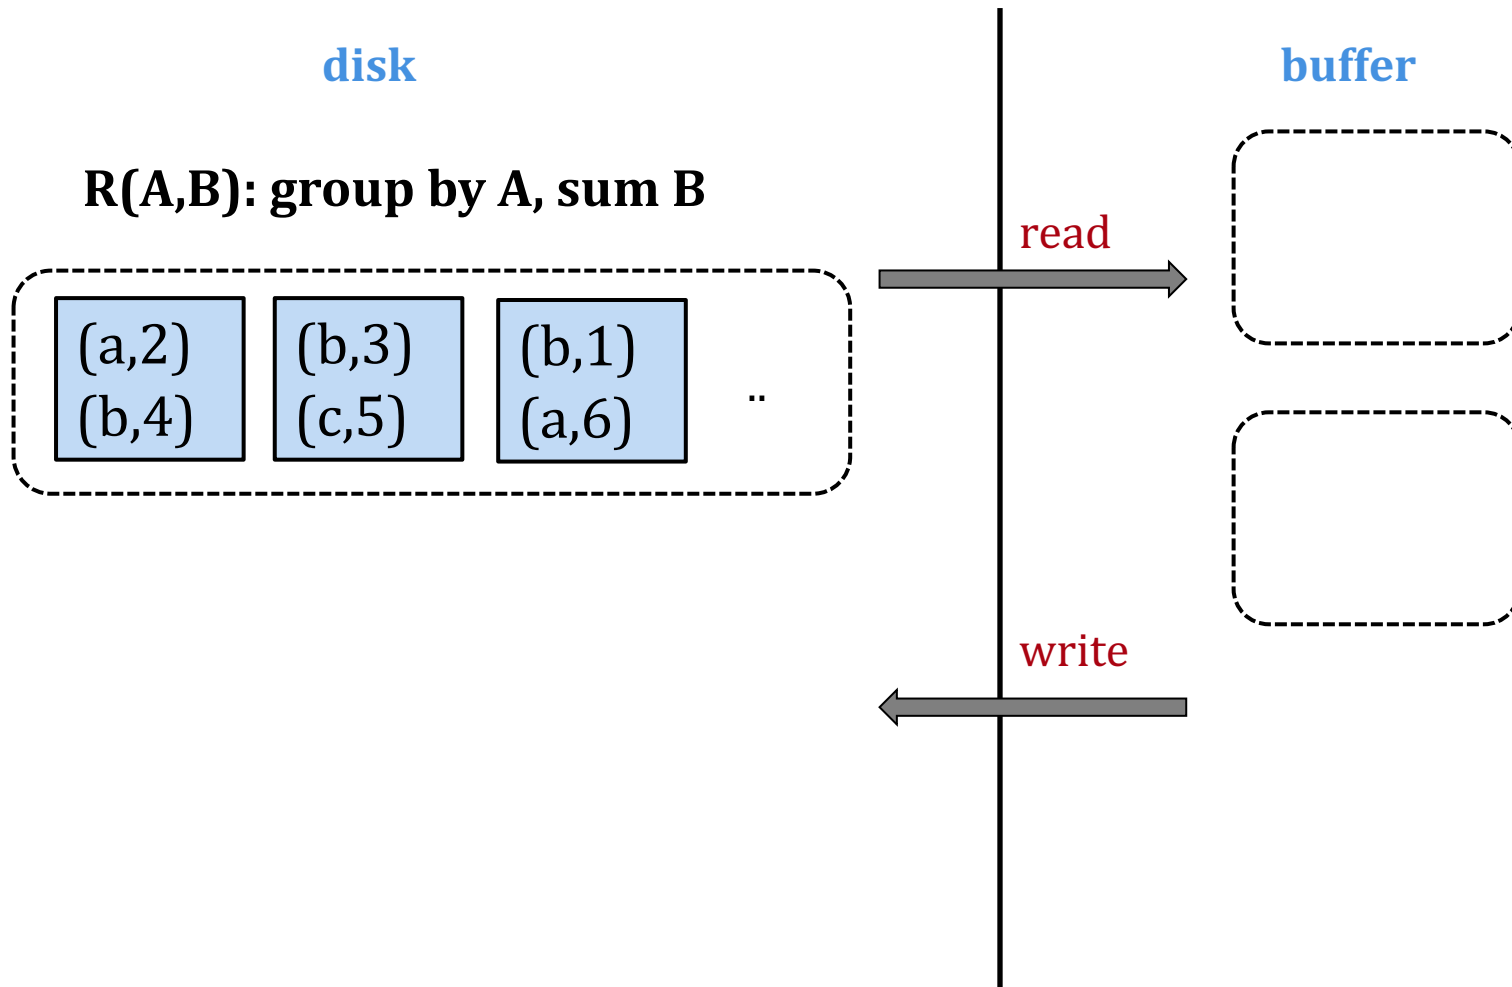
# AGGREGATION: SORTING

---

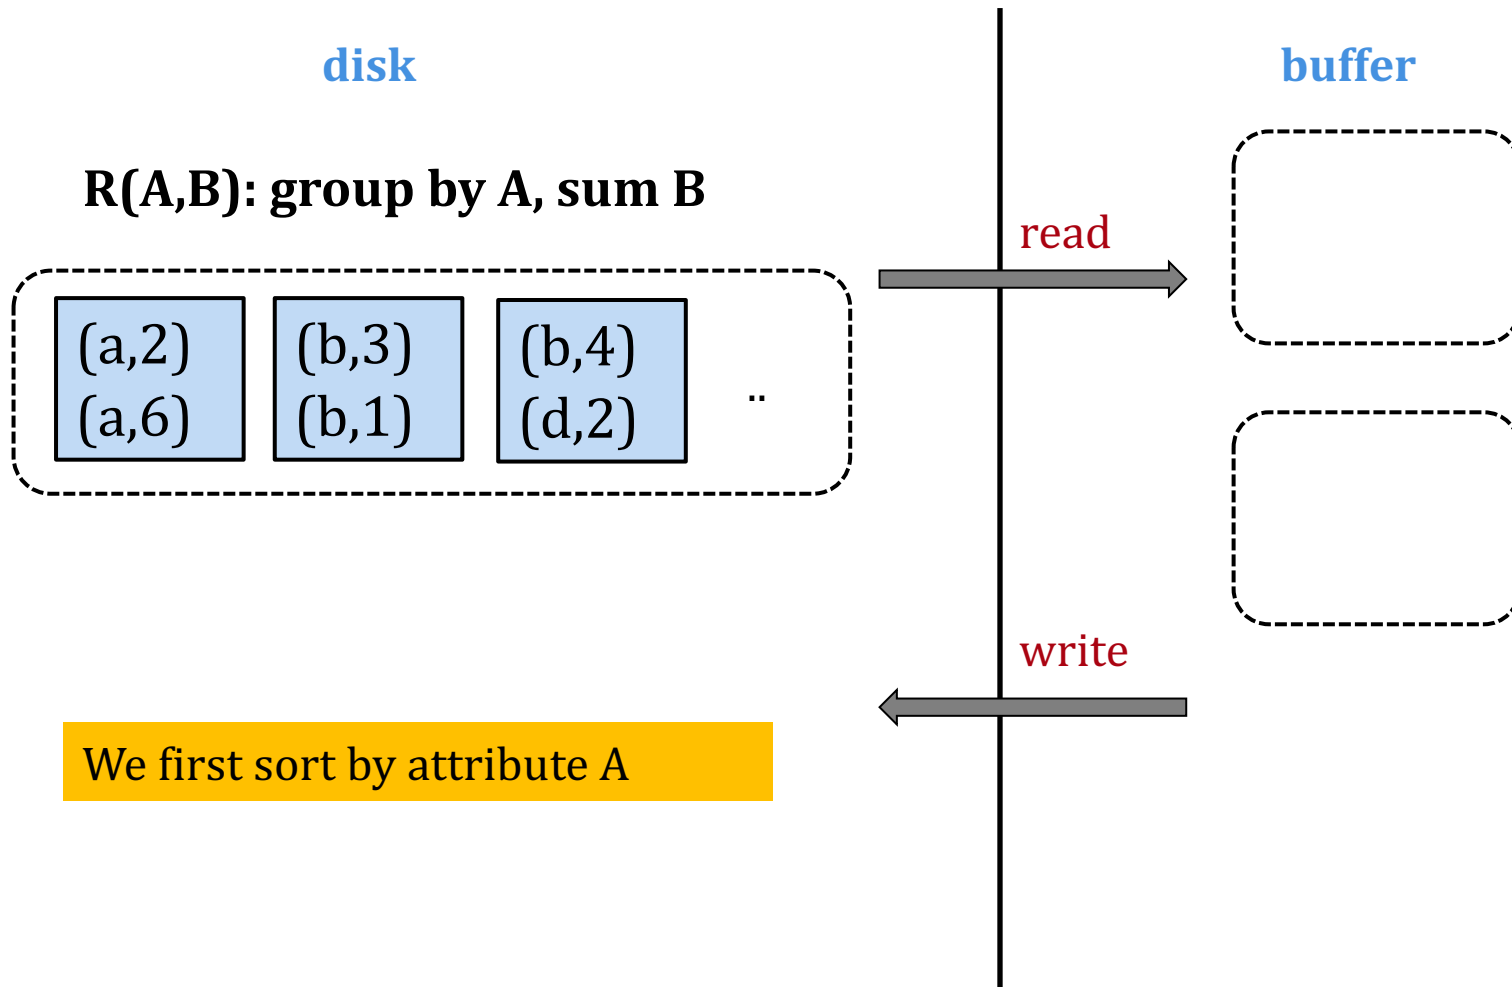
- sort on group by attributes (if any)
- scan sorted tuples, computing running aggregate
  - max/min: max/min
  - average: sum, count
- when the group by attribute changes, output aggregate result
- **cost** = sorting cost



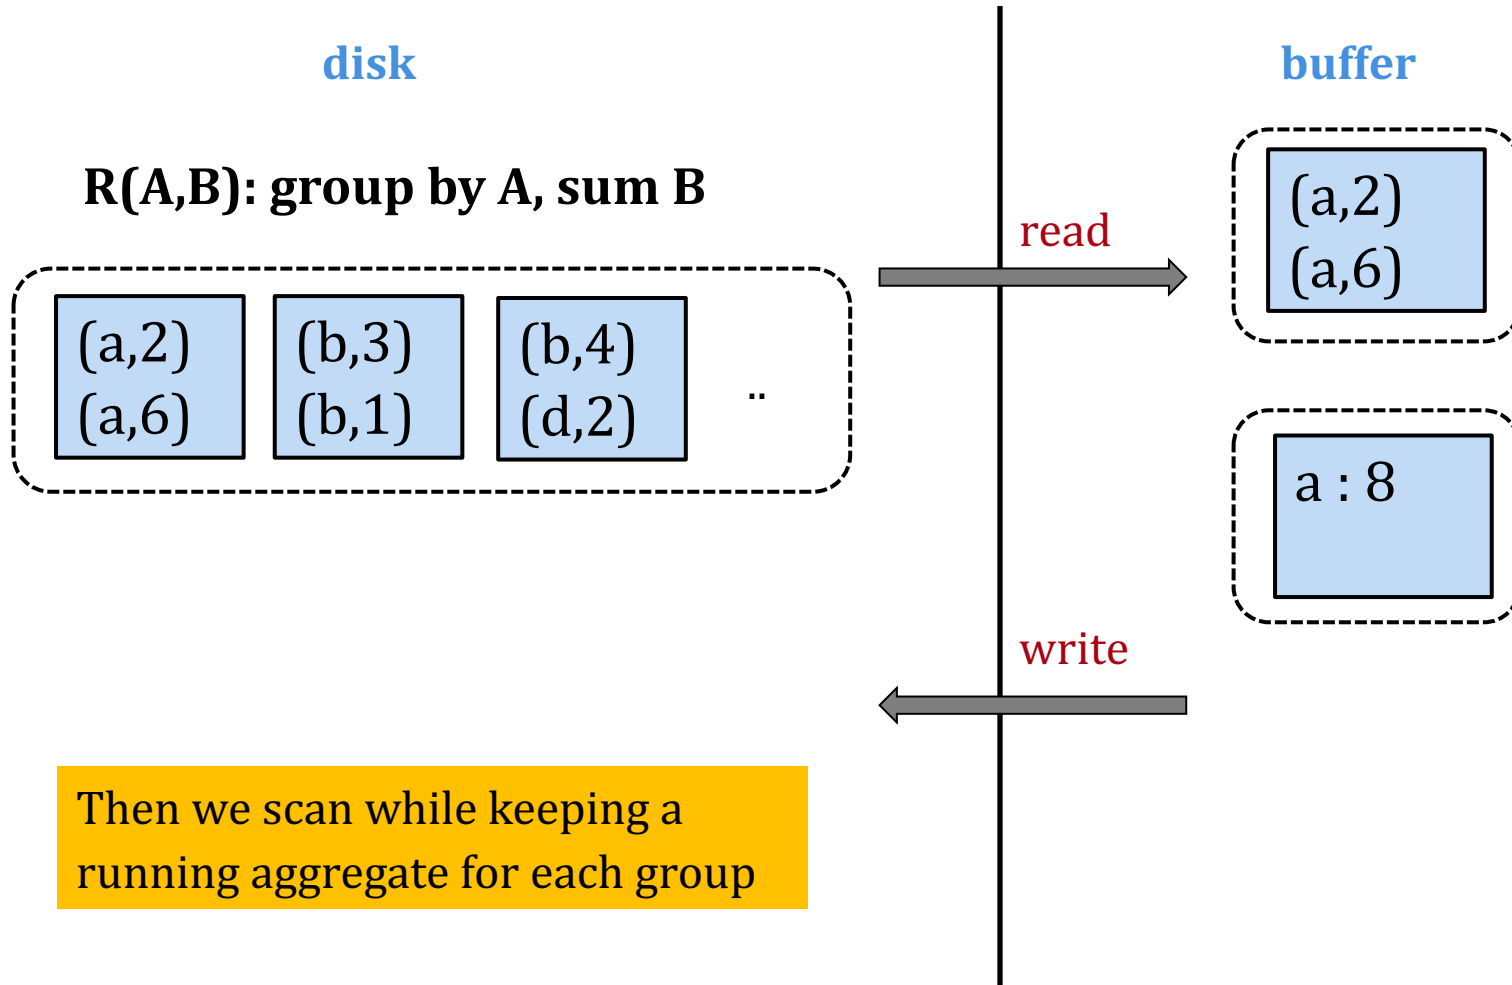
# EXAMPLE



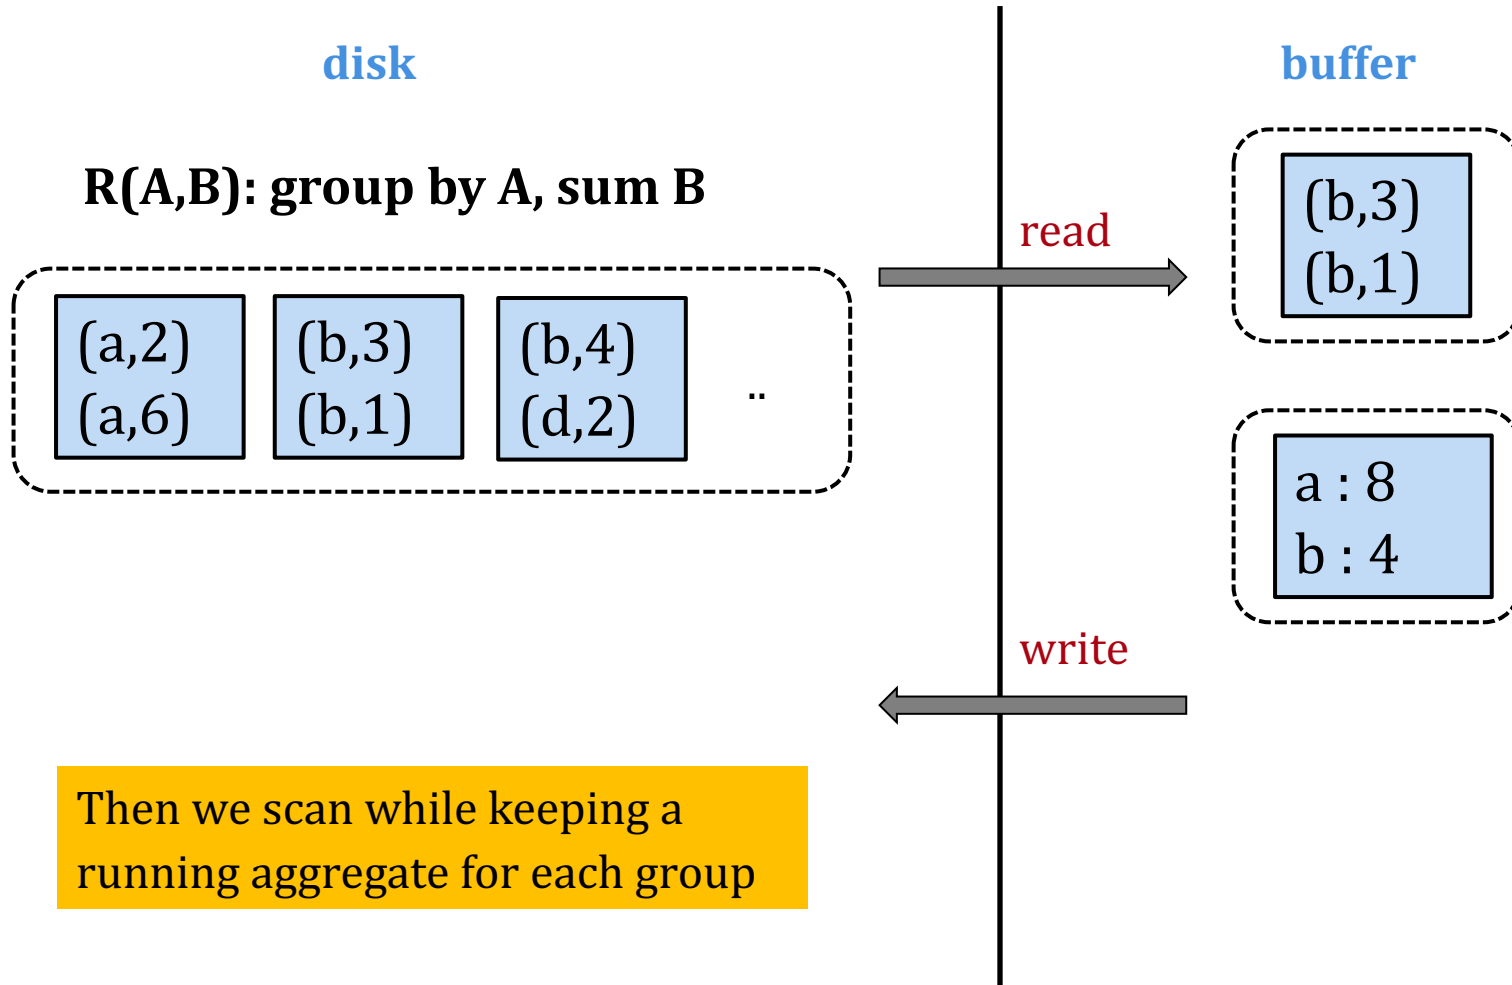
# EXAMPLE



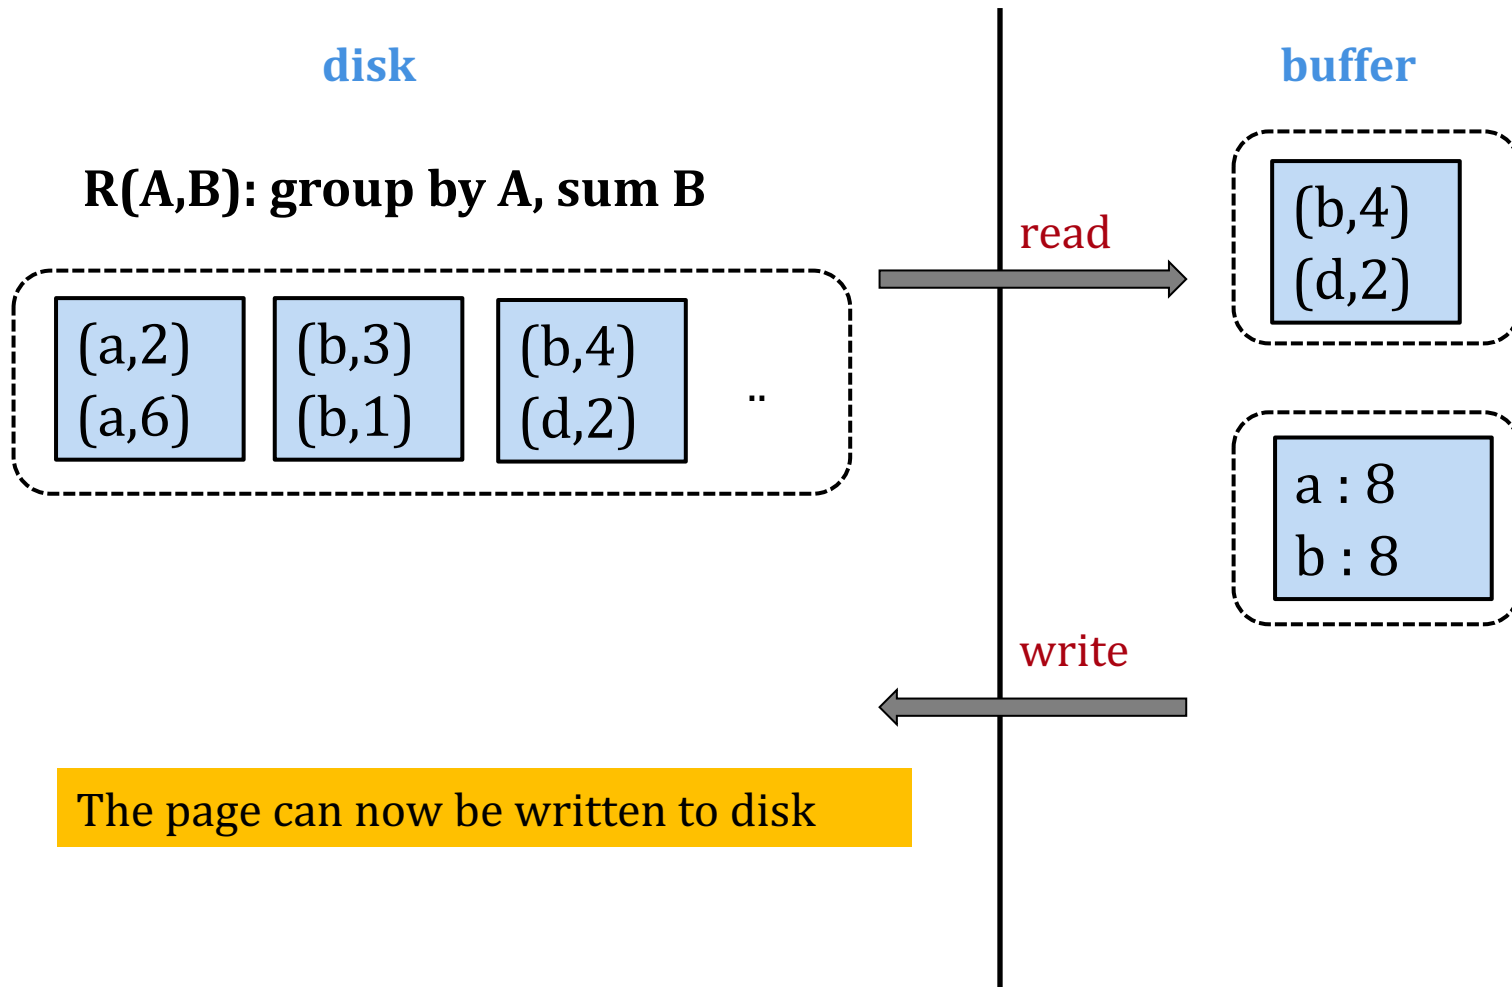
# EXAMPLE



# EXAMPLE



# EXAMPLE



The page can now be written to disk

---

# AGGREGATION: HASHING

---

- Hash on group by attributes (if any)
  - **Hash entry** = group attributes + running aggregate
- Scan tuples, probe hash table, update hash entry
- Scan hash table, and output each hash entry
- **cost** = scan relation

*What happens if we have too many groups?*