
CONTENTS

PREFACE	xxiv
Part I FOUNDATIONS	1
1 OVERVIEW OF DATABASE SYSTEMS	3
1.1 Managing Data	4
1.2 A Historical Perspective	6
1.3 File Systems versus a DBMS	8
1.4 Advantages of a DBMS	9
1.5 Describing and Storing Data in a DBMS	10
1.5.1 The Relational Model	11
1.5.2 Levels of Abstraction in a DBMS	12
1.5.3 Data Independence	15
1.6 Queries in a DBMS	16
1.7 Transaction Management	17
1.7.1 Concurrent Execution of Transactions	17
1.7.2 Incomplete Transactions and System Crashes	18
1.7.3 Points to Note	19
1.8 Structure of a DBMS	19
1.9 People Who Work with Databases	21
1.10 Review Questions	22
2 INTRODUCTION TO DATABASE DESIGN	25
2.1 Database Design and ER Diagrams	26
2.1.1 Beyond ER Design	27
2.2 Entities, Attributes, and Entity Sets	28
2.3 Relationships and Relationship Sets	29
2.4 Additional Features of the ER Model	32
2.4.1 Key Constraints	32
2.4.2 Participation Constraints	34
2.4.3 Weak Entities	35
2.4.4 Class Hierarchies	37
2.4.5 Aggregation	39

2.5	Conceptual Design With the ER Model	40
2.5.1	Entity versus Attribute	41
2.5.2	Entity versus Relationship	42
2.5.3	Binary versus Ternary Relationships	43
2.5.4	Aggregation versus Ternary Relationships	45
2.6	Conceptual Design for Large Enterprises	46
2.7	The Unified Modeling Language	47
2.8	Case Study: The Internet Shop	49
2.8.1	Requirements Analysis	49
2.8.2	Conceptual Design	50
2.9	Review Questions	51
3	THE RELATIONAL MODEL	57
3.1	Introduction to the Relational Model	59
3.1.1	Creating and Modifying Relations Using SQL	62
3.2	Integrity Constraints over Relations	63
3.2.1	Key Constraints	64
3.2.2	Foreign Key Constraints	66
3.2.3	General Constraints	68
3.3	Enforcing Integrity Constraints	69
3.3.1	Transactions and Constraints	72
3.4	Querying Relational Data	73
3.5	Logical Database Design: ER to Relational	74
3.5.1	Entity Sets to Tables	75
3.5.2	Relationship Sets (without Constraints) to Tables	76
3.5.3	Translating Relationship Sets with Key Constraints	78
3.5.4	Translating Relationship Sets with Participation Constraints	79
3.5.5	Translating Weak Entity Sets	82
3.5.6	Translating Class Hierarchies	83
3.5.7	Translating ER Diagrams with Aggregation	84
3.5.8	ER to Relational: Additional Examples	85
3.6	Introduction to Views	86
3.6.1	Views, Data Independence, Security	87
3.6.2	Updates on Views	88
3.7	Destroying/Altering Tables and Views	91
3.8	Case Study: The Internet Store	92
3.9	Review Questions	94
4	RELATIONAL ALGEBRA AND CALCULUS	100
4.1	Preliminaries	101
4.2	Relational Algebra	102
4.2.1	Selection and Projection	103
4.2.2	Set Operations	104

4.2.3	Renaming	106
4.2.4	Joins	107
4.2.5	Division	109
4.2.6	More Examples of Algebra Queries	110
4.3	Relational Calculus	116
4.3.1	Tuple Relational Calculus	117
4.3.2	Domain Relational Calculus	122
4.4	Expressive Power of Algebra and Calculus	124
4.5	Review Questions	126
5	SQL: QUERIES, CONSTRAINTS, TRIGGERS	130
5.1	Overview	131
5.1.1	Chapter Organization	132
5.2	The Form of a Basic SQL Query	133
5.2.1	Examples of Basic SQL Queries	138
5.2.2	Expressions and Strings in the SELECT Command	139
5.3	UNION, INTERSECT, and EXCEPT	141
5.4	Nested Queries	144
5.4.1	Introduction to Nested Queries	145
5.4.2	Correlated Nested Queries	147
5.4.3	Set-Comparison Operators	148
5.4.4	More Examples of Nested Queries	149
5.5	Aggregate Operators	151
5.5.1	The GROUP BY and HAVING Clauses	154
5.5.2	More Examples of Aggregate Queries	158
5.6	Null Values	162
5.6.1	Comparisons Using Null Values	163
5.6.2	Logical Connectives AND, OR, and NOT	163
5.6.3	Impact of SQL Constructs	163
5.6.4	Outer Joins	164
5.6.5	Disallowing Null Values	165
5.7	Complex Integrity Constraints in SQL	165
5.7.1	Constraints over a Single Table	165
5.7.2	Domain Constraints and Distinct Types	166
5.7.3	Assertions: ICs over Several Tables	167
5.8	Triggers and Active Databases	168
5.8.1	Examples of Triggers in SQL	169
5.9	Designing Active Databases	171
5.9.1	Why Triggers Can Be Hard to Understand	171
5.9.2	Constraints versus Triggers	172
5.9.3	Other Uses of Triggers	172
5.10	Review Questions	173

Part II	APPLICATION DEVELOPMENT	183
6	DATABASE APPLICATION DEVELOPMENT	185
6.1	Accessing Databases from Applications	187
6.1.1	Embedded SQL	187
6.1.2	Cursors	189
6.1.3	Dynamic SQL	194
6.2	An Introduction to JDBC	194
6.2.1	Architecture	196
6.3	JDBC Classes and Interfaces	197
6.3.1	JDBC Driver Management	197
6.3.2	Connections	198
6.3.3	Executing SQL Statements	200
6.3.4	ResultSets	201
6.3.5	Exceptions and Warnings	203
6.3.6	Examining Database Metadata	204
6.4	SQLJ	206
6.4.1	Writing SQLJ Code	207
6.5	Stored Procedures	209
6.5.1	Creating a Simple Stored Procedure	209
6.5.2	Calling Stored Procedures	210
6.5.3	SQL/PSM	212
6.6	Case Study: The Internet Book Shop	214
6.7	Review Questions	216
7	INTERNET APPLICATIONS	220
7.1	Introduction	220
7.2	Internet Concepts	221
7.2.1	Uniform Resource Identifiers	221
7.2.2	The Hypertext Transfer Protocol (HTTP)	223
7.3	HTML Documents	226
7.4	XML Documents	227
7.4.1	Introduction to XML	228
7.4.2	XML DTDs	231
7.4.3	Domain-Specific DTDs	234
7.5	The Three-Tier Application Architecture	236
7.5.1	Single-Tier and Client-Server Architectures	236
7.5.2	Three-Tier Architectures	239
7.5.3	Advantages of the Three-Tier Architecture	241
7.6	The Presentation Layer	242
7.6.1	HTML Forms	242
7.6.2	JavaScript	245
7.6.3	Style Sheets	247

7.7	The Middle Tier	251
7.7.1	CGI: The Common Gateway Interface	251
7.7.2	Application Servers	252
7.7.3	Servlets	254
7.7.4	JavaServer Pages	256
7.7.5	Maintaining State	258
7.8	Case Study: The Internet Book Shop	261
7.9	Review Questions	264
Part III	STORAGE AND INDEXING	271
8	OVERVIEW OF STORAGE AND INDEXING	273
8.1	Data on External Storage	274
8.2	File Organizations and Indexing	275
8.2.1	Clustered Indexes	277
8.2.2	Primary and Secondary Indexes	277
8.3	Index Data Structures	278
8.3.1	Hash-Based Indexing	279
8.3.2	Tree-Based Indexing	280
8.4	Comparison of File Organizations	282
8.4.1	Cost Model	283
8.4.2	Heap Files	284
8.4.3	Sorted Files	285
8.4.4	Clustered Files	287
8.4.5	Heap File with Unclustered Tree Index	288
8.4.6	Heap File With Unclustered Hash Index	289
8.4.7	Comparison of I/O Costs	290
8.5	Indexes and Performance Tuning	291
8.5.1	Impact of the Workload	292
8.5.2	Clustered Index Organization	292
8.5.3	Composite Search Keys	295
8.5.4	Index Specification in SQL:1999	299
8.6	Review Questions	299
9	STORING DATA: DISKS AND FILES	304
9.1	The Memory Hierarchy	305
9.1.1	Magnetic Disks	306
9.1.2	Performance Implications of Disk Structure	308
9.2	Redundant Arrays of Independent Disks	309
9.2.1	Data Striping	310
9.2.2	Redundancy	311
9.2.3	Levels of Redundancy	312
9.2.4	Choice of RAID Levels	316

9.3	Disk Space Management	316
9.3.1	Keeping Track of Free Blocks	317
9.3.2	Using as File Systems to Manage Disk Space	317
9.4	Buffer Manager	318
9.4.1	Buffer Replacement Policies	320
9.4.2	Buffer Management in DBMS versus OS	322
9.5	Files of Records	324
9.5.1	Implementing Heap Files	324
9.6	Page Formats	326
9.6.1	Fixed-Length Records	327
9.6.2	Variable-Length Records	328
9.7	Record Formats	330
9.7.1	Fixed-Length Records	331
9.7.2	Variable-Length Records	331
9.8	Review Questions	333
10	TREE-STRUCTURED INDEXING	338
10.1	Intuition For Tree Indexes	339
10.2	Indexed Sequential Access Method (ISAM)	341
10.2.1	Overflow Pages, Locking Considerations	344
10.3	B+ Trees: A Dynamic Index Structure	344
10.3.1	Format of a Node	346
10.4	Search	347
10.5	Insert	348
10.6	Delete	352
10.7	Duplicates	356
10.8	B+ Trees in Practice	358
10.8.1	Key Compression	358
10.8.2	Bulk-Loading a B+ Tree	360
10.8.3	The Order Concept	363
10.8.4	The Effect of Inserts and Deletes on Rids	364
10.9	Review Questions	364
11	HASH-BASED INDEXING	370
11.1	Static Hashing	371
11.1.1	Notation and Conventions	373
11.2	Extendible Hashing	373
11.3	Linear Hashing	379
11.4	Extendible vs. Linear Hashing	384
11.5	Review Questions	385
Part IV	QUERY EVALUATION	391

Contents

12	OVERVIEW OF QUERY EVALUATION	393
12.1	The System Catalog	394
12.1.1	Information in the Catalog	395
12.2	Introduction to Operator Evaluation	397
12.2.1	Three Common Techniques	398
12.2.2	Access Paths	398
12.3	Algorithms for Relational Operations	400
12.3.1	Selection	401
12.3.2	Projection	401
12.3.3	Join	402
12.3.4	Other Operations	404
12.4	Introduction to Query Optimization	404
12.4.1	Query Evaluation Plans	405
12.4.2	Multi-operator Queries: Pipelined Evaluation	407
12.4.3	The Iterator Interface	408
12.5	Alternative Plans: A Motivating Example	409
12.5.1	Pushing Selections	409
12.5.2	Using Indexes	411
12.6	What a Typical Optimizer Does	414
12.6.1	Alternative Plans Considered	414
12.6.2	Estimating the Cost of a Plan	416
12.7	Review Questions	417
13	EXTERNAL SORTING	421
13.1	When Does a DBMS Sort Data?	422
13.2	A Simple Two-Way Merge Sort	423
13.3	External Merge Sort	424
13.3.1	Minimizing the Number of Runs	428
13.4	Minimizing I/O Cost versus Number of I/Os	430
13.4.1	Blocked I/O	430
13.4.2	Double Buffering	432
13.5	Using B+ Trees for Sorting	433
13.5.1	Clustered Index	433
13.5.2	Unclustered Index	434
13.6	Review Questions	436
14	EVALUATING RELATIONAL OPERATORS	439
14.1	The Selection Operation	441
14.1.1	No Index, Unsorted Data	441
14.1.2	No Index, Sorted Data	442
14.1.3	B+ Tree Index	442
14.1.4	Hash Index, Equality Selection	444
14.2	General Selection Conditions	444

14.2.1	CNF and Index Matching	445
14.2.2	Evaluating Selections without Disjunction	445
14.2.3	Selections with Disjunction	446
14.3	The Projection Operation	447
14.3.1	Projection Based on Sorting	448
14.3.2	Projection Based on Hashing	449
14.3.3	Sorting Versus Hashing for Projections	451
14.3.4	Use of Indexes for Projections	452
14.4	The Join Operation	452
14.4.1	Nested Loops Join	454
14.4.2	Sort-Merge Join	458
14.4.3	Hash Join	463
14.4.4	General Join Conditions	467
14.5	The Set Operations	468
14.5.1	Sorting for Union and Difference	469
14.5.2	Hashing for Union and Difference	469
14.6	Aggregate Operations	469
14.6.1	Implementing Aggregation by Using an Index	471
14.7	The Impact of Buffering	471
14.8	Review Questions	472
15	A TYPICAL RELATIONAL QUERY OPTIMIZER	478
15.1	Translating SQL Queries into Algebra	479
15.1.1	Decomposition of a Query into Blocks	479
15.1.2	A Query Block as a Relational Algebra Expression	481
15.2	Estimating the Cost of a Plan	482
15.2.1	Estimating Result Sizes	483
15.3	Relational Algebra Equivalences	488
15.3.1	Selections	488
15.3.2	Projections	488
15.3.3	Cross-Products and Joins	489
15.3.4	Selects, Projects, and Joins	490
15.3.5	Other Equivalences	491
15.4	Enumeration of Alternative Plans	492
15.4.1	Single-Relation Queries	492
15.4.2	Multiple-Relation Queries	496
15.5	Nested Subqueries	504
15.6	The System R Optimizer	506
15.7	Other Approaches to Query Optimization	507
15.8	Review Questions	507
Part V	TRANSACTION MANAGEMENT	517

16	OVERVIEW OF TRANSACTION MANAGEMENT	519
16.1	The ACID Properties	520
16.1.1	Consistency and Isolation	521
16.1.2	Atomicity and Durability	522
16.2	Transactions and Schedules	523
16.3	Concurrent Execution of Transactions	524
16.3.1	Motivation for Concurrent Execution	524
16.3.2	Serializability	525
16.3.3	Anomalies Due to Interleaved Execution	526
16.3.4	Schedules Involving Aborted Transactions	529
16.4	Lock-Based Concurrency Control	530
16.4.1	Strict Two-Phase Locking (Strict 2PL)	531
16.4.2	Deadlocks	533
16.5	Performance of Locking	533
16.6	Transaction Support in SQL	535
16.6.1	Creating and Terminating Transactions	535
16.6.2	What Should We Lock?	537
16.6.3	Transaction Characteristics in SQL	538
16.7	Introduction to Crash Recovery	540
16.7.1	Stealing Frames and Forcing Pages	541
16.7.2	Recovery-Related Steps during Normal Execution	542
16.7.3	Overview of ARIES	543
16.7.4	Atomicity: Implementing Rollback	543
16.8	Review Questions	544
17	CONCURRENCY CONTROL	549
17.1	2PL, Serializability, and Recoverability	550
17.1.1	View Serializability	553
17.2	Introduction to Lock Management	553
17.2.1	Implementing Lock and Unlock Requests	554
17.3	Lock Conversions	555
17.4	Dealing With Deadlocks	556
17.4.1	Deadlock Prevention	558
17.5	Specialized Locking Techniques	559
17.5.1	Dynamic Databases and the Phantom Problem	560
17.5.2	Concurrency Control in B+ Trees	561
17.5.3	Multiple-Granularity Locking	564
17.6	Concurrency Control without Locking	566
17.6.1	Optimistic Concurrency Control	566
17.6.2	Timestamp-Based Concurrency Control	569
17.6.3	Multiversion Concurrency Control	572
17.7	Review Questions	573

18	CRASH RECOVERY	579
18.1	Introduction to ARIES	580
18.2	The Log	582
18.3	Other Recovery-Related Structures	585
18.4	The Write-Ahead Log Protocol	586
18.5	Checkpointing	587
18.6	Recovering from a System Crash	587
	18.6.1 Analysis Phase	588
	18.6.2 Redo Phase	590
	18.6.3 Undo Phase	592
18.7	Media Recovery	595
18.8	Other Approaches and Interaction with Concurrency Control	596
18.9	Review Questions	597
Part VI	DATABASE DESIGN AND TUNING	603
19	SCHEMA REFINEMENT AND NORMAL FORMS	605
19.1	Introduction to Schema Refinement	606
	19.1.1 Problems Caused by Redundancy	606
	19.1.2 Decompositions	608
	19.1.3 Problems Related to Decomposition	609
19.2	Functional Dependencies	611
19.3	Reasoning about FDs	612
	19.3.1 Closure of a Set of FDs	612
	19.3.2 Attribute Closure	614
19.4	Normal Forms	615
	19.4.1 Boyce-Codd Normal Form	615
	19.4.2 Third Normal Form	617
19.5	Properties of Decompositions	619
	19.5.1 Lossless-Join Decomposition	619
	19.5.2 Dependency-Preserving Decomposition	621
19.6	Normalization	622
	19.6.1 Decomposition into BCNF	622
	19.6.2 Decomposition into 3NF	625
19.7	Schema Refinement in Database Design	629
	19.7.1 Constraints on an Entity Set	630
	19.7.2 Constraints on a Relationship Set	630
	19.7.3 Identifying Attributes of Entities	631
	19.7.4 Identifying Entity Sets	633
19.8	Other Kinds of Dependencies	633
	19.8.1 Multivalued Dependencies	634
	19.8.2 Fourth Normal Form	636
	19.8.3 Join Dependencies	638

19.8.4	Fifth Normal Form	638
19.8.5	Inclusion Dependencies	639
19.9	Case Study: The Internet Shop	640
19.10	Review Questions	642
20	PHYSICAL DATABASE DESIGN AND TUNING	649
20.1	Introduction to Physical Database Design	650
20.1.1	Database Workloads	651
20.1.2	Physical Design and Tuning Decisions	652
20.1.3	Need for Database Tuning	653
20.2	Guidelines for Index Selection	653
20.3	Basic Examples of Index Selection	656
20.4	Clustering and Indexing	658
20.4.1	Co-clustering Two Relations	660
20.5	Indexes that Enable Index-Only Plans	662
20.6	Tools to Assist in Index Selection	663
20.6.1	Automatic Index Selection	663
20.6.2	How Do Index Tuning Wizards Work?	664
20.7	Overview of Database Tuning	667
20.7.1	Tuning Indexes	667
20.7.2	Tuning the Conceptual Schema	669
20.7.3	Tuning Queries and Views	670
20.8	Choices in Tuning the Conceptual Schema	671
20.8.1	Settling for a Weaker Normal Form	671
20.8.2	Denormalization	672
20.8.3	Choice of Decomposition	672
20.8.4	Vertical Partitioning of BCNF Relations	674
20.8.5	Horizontal Decomposition	674
20.9	Choices in Tuning Queries and Views	675
20.10	Impact of Concurrency	678
20.10.1	Reducing Lock Durations	678
20.10.2	Reducing Hot Spots	679
20.11	Case Study: The Internet Shop	680
20.11.1	Tuning the Database	682
20.12	DBMS Benchmarking	682
20.12.1	Well-Known DBMS Benchmarks	683
20.12.2	Using a Benchmark	684
20.13	Review Questions	685
21	SECURITY AND AUTHORIZATION	692
21.1	Introduction to Database Security	693
21.2	Access Control	694
21.3	Discretionary Access Control	695

21.3.1	Grant and Revoke on Views and Integrity Constraints	704
21.4	Mandatory Access Control	705
21.4.1	Multilevel Relations and Polyinstantiation	707
21.4.2	Covert Channels, DoD Security Levels	708
21.5	Security for Internet Applications	709
21.5.1	Encryption	709
21.5.2	Certifying Servers: The SSL Protocol	712
21.5.3	Digital Signatures	713
21.6	Additional Issues Related to Security	714
21.6.1	Role of the Database Administrator	714
21.6.2	Security in Statistical Databases	715
21.7	Design Case Study: The Internet Store	716
21.8	Review Questions	718
Part VII	ADDITIONAL TOPICS	723
22	PARALLEL AND DISTRIBUTED DATABASES	725
22.1	Introduction	726
22.2	Architectures for Parallel Databases	727
22.3	Parallel Query Evaluation	728
22.3.1	Data Partitioning	730
22.3.2	Parallelizing Sequential Operator Evaluation Code	730
22.4	Parallelizing Individual Operations	731
22.4.1	Bulk Loading and Scanning	731
22.4.2	Sorting	732
22.4.3	Joins	732
22.5	Parallel Query Optimization	735
22.6	Introduction to Distributed Databases	736
22.6.1	Types of Distributed Databases	737
22.7	Distributed DBMS Architectures	737
22.7.1	Client-Server Systems	738
22.7.2	Collaborating Server Systems	738
22.7.3	Middleware Systems	739
22.8	Storing Data in a Distributed DBMS	739
22.8.1	Fragmentation	739
22.8.2	Replication	741
22.9	Distributed Catalog Management	741
22.9.1	Naming Objects	741
22.9.2	Catalog Structure	742
22.9.3	Distributed Data Independence	743
22.10	Distributed Query Processing	743
22.10.1	Nonjoin Queries in a Distributed DBMS	744
22.10.2	Joins in a Distributed DBMS	745

22.10.3 Cost-Based Query Optimization	749
22.11 Updating Distributed Data	750
22.11.1 Synchronous Replication	750
22.11.2 Asynchronous Replication	751
22.12 Distributed Transactions	755
22.13 Distributed Concurrency Control	755
22.13.1 Distributed Deadlock	756
22.14 Distributed Recovery	758
22.14.1 Normal Execution and Commit Protocols	758
22.14.2 Restart after a Failure	760
22.14.3 Two-Phase Commit Revisited	761
22.14.4 Three-Phase Commit	762
22.15 Review Questions	763
23 OBJECT-DATABASE SYSTEMS	772
23.1 Motivating Example	774
23.1.1 New Data Types	775
23.1.2 Manipulating the New Data	777
23.2 Structured Data Types	779
23.2.1 Collection Types	780
23.3 Operations on Structured Data	781
23.3.1 Operations on Rows	781
23.3.2 Operations on Arrays	781
23.3.3 Operations on Other Collection Types	782
23.3.4 Queries Over Nested Collections	783
23.4 Encapsulation and ADTs	784
23.4.1 Defining Methods	785
23.5 Inheritance	787
23.5.1 Defining Types with Inheritance	787
23.5.2 Binding Methods	788
23.5.3 Collection Hierarchies	789
23.6 Objects, aIDs, and Reference Types	789
23.6.1 Notions of Equality	790
23.6.2 Dereferencing Reference Types	791
23.6.3 URLs and OIDs in SQL:1999	791
23.7 Database Design for an ORDBMS	792
23.7.1 Collection Types and ADTs	792
23.7.2 Object Identity	795
23.7.3 Extending the ER Model	796
23.7.4 Using Nested Collections	798
23.8 ORDBMS Implementation Challenges	799
23.8.1 Storage and Access Methods	799
23.8.2 Query Processing	801

DATABASE MANAGEMENT SYSTEMS

23.8.3	Query Optimization	803
23.9	OODBMS	805
23.9.1	The ODMG Data Model and ODL	805
23.9.2	OQL	807
23.10	Comparing RDBMS, OODBMS, and ORDBMS	809
23.10.1	RDBMS versus ORDBMS	809
23.10.2	OODBMS versus ORDBMS: Similarities	809
23.10.3	OODBMS versus ORDBMS: Differences	810
23.11	Review Questions	811
24	DEDUCTIVE DATABASES	817
24.1	Introduction to Recursive Queries	818
24.1.1	Datalog	819
24.2	Theoretical Foundations	822
24.2.1	Least Model Semantics	823
24.2.2	The Fixpoint Operator	824
24.2.3	Safe Datalog Programs	825
24.2.4	Least Model = Least Fixpoint	826
24.3	Recursive Queries with Negation	827
24.3.1	Stratification	828
24.4	From Datalog to SQL	831
24.5	Evaluating Recursive Queries	834
24.5.1	Fixpoint Evaluation without Repeated Inferences	835
24.5.2	Pushing Selections to Avoid Irrelevant Inferences	837
24.5.3	The Magic Sets Algorithm	838
24.6	Review Questions	841
25	DATA WAREHOUSING AND DECISION SUPPORT	846
25.1	Introduction to Decision Support	848
25.2	OLAP: Multidimensional Data Model	849
25.2.1	Multidimensional Database Design	853
25.3	Multidimensional Aggregation Queries	854
25.3.1	ROLLUP and CUBE in SQL:1999	856
25.4	Window Queries in SQL:1999	859
25.4.1	Framing a Window	861
25.4.2	New Aggregate Functions	862
25.5	Findipg Answers Quickly	862
25.5.1	Top N Queries	863
25.5.2	Online Aggregation	864
25.6	Implementation Techniques for OLAP	865
25.6.1	Bitmap Indexes	866
25.6.2	Join Indexes	868
25.6.3	File Organizations	869

Contents

25.7	Data Warehousing	870
25.7.1	Creating and Maintaining a Warehouse	871
25.8	Views and Decision Support	872
25.8.1	Views, OLAP, and Warehousing	872
25.8.2	Queries over Views	873
25.9	View Materialization	873
25.9.1	Issues in View Materialization	874
25.10	Maintaining Materialized Views	876
25.10.1	Incremental View Maintenance	876
25.10.2	Maintaining Warehouse Views	879
25.10.3	When Should We Synchronize Views?	881
25.11	Review Questions	882
26	DATA MINING	889
26.1	Introduction to Data Mining	890
26.1.1	The Knowledge Discovery Process	891
26.2	Counting Co-occurrences	892
26.2.1	Frequent Itemsets	892
26.2.2	Iceberg Queries	895
26.3	Mining for Rules	897
26.3.1	Association Rules	897
26.3.2	An Algorithm for Finding Association Rules	898
26.3.3	Association Rules and ISA Hierarchies	899
26.3.4	Generalized Association Rules	900
26.3.5	Sequential Patterns	901
26.3.6	The Use of Association Rules for Prediction	902
26.3.7	Bayesian Networks	903
26.3.8	Classification and Regression Rules	904
26.4	Tree-Structured Rules	906
26.4.1	Decision Trees	907
26.4.2	An Algorithm to Build Decision Trees	908
26.5	Clustering	911
26.5.1	A Clustering Algorithm	912
26.6	Similarity Search over Sequences	913
26.6.1	An Algorithm to Find Similar Sequences	915
26.7	Incremental Mining and Data Streams	916
26.7.1	Incremental Maintenance of Frequent Itemsets	918
26.8	Additional Data Mining Tasks	920
26.9	Review Questions	920
27	INFORMATION RETRIEVAL AND XML DATA	926
27.1	Colliding Worlds: Databases, IR, and XML	927
27.1.1	DBMS versus IR Systems	928

27.2	Introduction to Information Retrieval	929
27.2.1	Vector Space Model	930
27.2.2	TF <i>idf</i> Weighting of Terms	931
27.2.3	Ranking Document Similarity	932
27.2.4	Measuring Success: Precision and Recall	934
27.3	Indexing for Text Search	934
27.3.1	Inverted Indexes	935
27.3.2	Signature Files	937
27.4	Web Search Engines	939
27.4.1	Search Engine Architecture	939
27.4.2	Using Link Information	940
27.5	Managing Text in a DBMS	944
27.5.1	Loosely Coupled Inverted Index	945
27.6	A Data Model for XML	945
27.6.1	Motivation for Loose Structure	945
27.6.2	A Graph Model	946
27.7	XQuery: Querying XML Data	948
27.7.1	Path Expressions	948
27.7.2	FLWR Expressions	949
27.7.3	Ordering of Elements	951
27.7.4	Grouping and Generation of Collection Values	951
27.8	Efficient Evaluation of XML Queries	952
27.8.1	Storing XML in RDBMS	952
27.8.2	Indexing XML Repositories	956
27.9	Review Questions	959
28	SPATIAL DATA MANAGEMENT	968
28.1	Types of Spatial Data and Queries	969
28.2	Applications Involving Spatial Data	971
28.3	Introduction to Spatial Indexes	973
28.3.1	Overview of Proposed Index Structures	974
28.4	Indexing Based on Space-Filling Curves	975
28.4.1	Region Quad Trees and Z-Ordering: Region Data	976
28.4.2	Spatial Queries Using Z-Ordering	978
28.5	Grid Files	978
28.5.1	Adapting Grid Files to Handle Regions	981
28.6	R Trees: Point and Region Data	982
28.6.1	Queries	983
28.6.2	Insert and Delete Operations	984
28.6.3	Concurrency Control	986
28.6.4	Generalized Search Trees	987
28.7	Issues in High-Dimensional Indexing	988
28.8	Review Questions	988

29 FURTHER READING	992
29.1 Advanced Transaction Processing	993
29.1.1 Transaction Processing Monitors	993
29.1.2 New Transaction Models	994
29.1.3 Real-Time DBMSs	994
29.2 Data Integration	995
29.3 Mobile Databases	995
29.4 Main Memory Databases	996
29.5 Multimedia Databases	997
29.6 Geographic Information Systems	998
29.7 Temporal Databases	999
29.8 Biological Databases	999
29.9 Information Visualization	1000
29.10 Summary	1000
30 THE MINIBASE SOFTWARE	1002
30.1 What Is Available	1002
30.2 Overview of Minibase Assignments	1003
30.3 Acknowledgments	1004
REFERENCES	1005
AUTHOR INDEX	1045
SUBJECT INDEX	1054